

# Analyzing the Ethereum Blockchain

Christin Sauerbier  
Humboldt University Berlin  
Berlin, Germany  
christin.sauerbier@outlook.com

Roman Proskalovich  
Humboldt University Berlin  
Berlin, Germany  
romanarion@gmail.com

Thomas Siskos  
Humboldt University Berlin  
Berlin, Germany  
thomas.siskos@hu-berlin.de

**Abstract**—In the explorative study we investigate the Ethereum blockchain, a decentralized crypto technology and a platform used for decentralized applications and currency transactions. The analysis covered 3 months of data, including transaction and economic data scraped from the Ethereum platform. The objective is to get closer insights into the Ethereum economy and find patterns in the data. The study reveals information about the main senders and receivers during the analyzed period and the underlying business processes, plus information about the network design. Furthermore, we scrutinize the period during which the hard fork of the Ethereum blockchain occurred. For the graph analysis the so called *Small World Phenomenon* could not be confirmed for the whole network, but for subgraphs it holds true.

**Index Terms**—ethereum, blockchain, cryptocurrencies, network analysis, descriptive analysis, graph analysis, crypto technology

## I. INTRODUCTION

With the Introduction of Bitcoin in 2008, a significant new currency was brought into life. It is based on a peer to peer network which gives the opportunity to cancel out intermediaries within money transfers, through a decentralized verification process. With that knowledge the idea was born to use such a technology to also run decentralized applications which can execute whole programs on such a platform. Ethereum can provide that and is now the second largest cryptocurrency after Bitcoin in terms of market capitalization [BitInfoCharts, 2019b], [BitInfoCharts, 2019a]. This paper proposes an explorative study on the Ethereum blockchain. The idea is to replicate what Lischke and Fabian [Lischke and Fabian, 2016] did and transfer it to another crypto technology the Ethereum Blockchain. They used the address spaces and transaction history of the Bitcoin and added off-set network data to get more insights about the bitcoin economy. The analyzed period included the first four years of the bitcoin existence. They were able to combine off-set network data and transaction data to get good statements about the business categories the currency is used for. With the Ethereum we try to figure out if we are able to collect the transaction data and IP addresses to get more insights about the crypto technology. Specifically, we want to figure out how much information one is able to get out of the transaction and off-set data to get deeper insights about what is going on in the network. Furthermore, we also want to assess if the information discloses security problems for the network. The

period is shorter than in the replicated paper, but we do used the same resources.

This paper is structured as follows. The first part is dedicated to the basics of Ethereum and how it got evolved from the bitcoin currency, or moreover the technology behind it. As we replicate a paper about the bitcoin blockchain [Lischke and Fabian, 2016] we also compare the Ethereum blockchain with the bitcoin blockchain. Give some relevant economic data about both. This paper is structured as follows. Section II and III are dedicated to the basics of Ethereum and how it got evolved from the bitcoin currency, as well as the technology behind it, respectively. As we replicate a paper about the bitcoin blockchain [Fabian, 2016] we also compare the Ethereum blockchain with the bitcoin blockchain. We follow this by a brief literature review in section IV. Afterwards we follow by demonstrating the underlying methodology in section V. The focus here lies mainly on how we obtained our dataset in subsection V-A and an overview of the methods used during graph analysis in subsection V-C. Moving onward, we present our findings in section VI, where we describe the outcome of our methods for data acquisition in subsection VI-A, structure them using traditional statistical methods in subsection VI-B and graph analysis in subsection VI-C. Finally, section VII concludes the article and discusses the findings of our analysis. We also want to use this part to discuss the main differences to the former paper about the bitcoin blockchain and give an idea about future work.

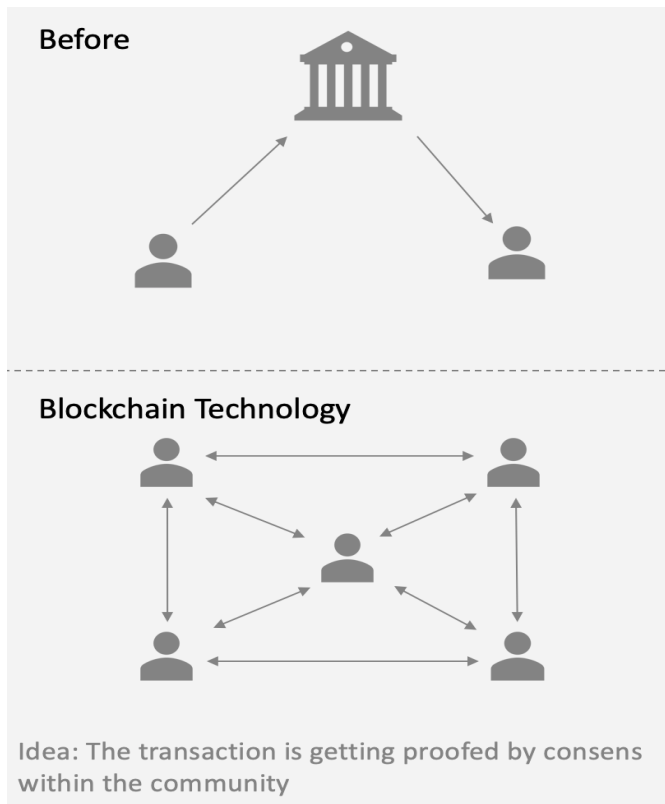
## II. ETHEREUM THEORY

Before diving into Ethereum one has to give a short introduction about the main invention of cryptocurrencies, Bitcoin and the blockchain technology behind it. These ideas drove Vitalik Buterin when he tried to design another cryptocurrency to take the bitcoin technology to the next level by giving developers to opportunity to write their own programs on the blockchain technology. This created the opportunity to use the blockchain technology and combine it with several business cases, not just as a currency.

The Bitcoin is a decentralized cryptocurrency which was invented by Satoshi Nakamoto, it combines three technologies with each other, Cryptography, Proof of Work and decentralized Networks. The idea was to create a currency which is independent from other institutions, who control the network as an intermediary. When using cryptocurrency, the network is

decentralized, and an intermediary becomes obsolete. The network takes over the function of an intermediary, as everyone keeps a transaction history, and this gets verified against each other, every time someone makes a new entry. This results in a technology that is hard to manipulate, meaning that it becomes increasingly hard for someone to take over, control or shut down the entire network [Grishchenko et al., 2018c].

Fig. 1: Blockchain technology is the basis for Ethereum




They could realize such a network with three already existing technologies which got combined the first time by Satoshi Nakamoto, for the Bitcoin. It describes the whole process how the Bitcoin blockchain gets assembled, validated and agreed for a transaction. The blockchain can be assumed as a database that is stored on several computers, in the same version [Grishchenko et al., 2018c].

The blockchain component can be described as a series of blocks that are combined ordered or chained with each other. Everyone stores a variety of information with a hash function. Were the information in every hash is related to the former one, plus the verification that the transaction is proofed. So far, most crypto technologies work with the so-called proof-of-work [Ray, 2018].

The main idea is to find a consent in the network which transactions or blocks are getting approved through the network. The work part in proof of work relates to the puzzle the community has to solve to verify a transaction and if enough transactions get combined, they can get inserted into a block.

This means that transactions can only be verified, if they are

Fig. 2: How Blocks get chained



gathered and inserted into a newly created block. The person in the network who is creating the new block is called miner and gets compensated with a value for mining. The exact compensation depends on the crypto technology. In Bitcoin if they verify for the Bitcoin network and in gas if they are mining for the Ethereum Blockchain. Every transaction gets collected through its hash and approved in a block which line up to a chain, so that it gets computationally hard to manipulate a transaction. Here one can lead over to the next important component of the bitcoin technology, the peer to peer networks.

Transferred it means, that there is a network in which each peer is connected to each other without any other intermediary in between. There is a direct connection between the participants within the network. It also is decentralized, such that nobody can manipulate anything.

When the creator of Ethereum thought about this idea and technology, the main question was, what else can be decentralized to cancel trusted intermediaries out and go beyond currencies. This let Vitalik Buterin in his white paper in late 2013, to describe a platform which can run decentralized applications and therefore run economic systems in pure software. The Vision of Ethereum is to create an unstoppable censorship-resistant self-sustaining decentralized world computer. [Ray, 2018]. The Ethereum can be explained best, if the dimensions get differentiated into the currency and the platform. The platform includes the Ethereum Virtual Machine, the smart contracts and the decentralized applications.

To add on to the former Bitcoin explanations, the Ethereum is also providing a virtual machine to realize smart contract code. Before diving into that topic it should be noted that the Ethereum platform is able to access and alter data to compute anything a Turing complete machine could, it is also able to use conditional branches. The last requirement to fulfill the Turing completeness is to have an arbitrary amount of memory. As not all requirements are met the Ethereum virtual machine (EVM) is only quasi-Turing complete, because every transaction or computation is requiring a defined amount of gas, that is used to pay the system for mining the blocks. The programming language to write applications for the EVM is Solidity, which gets compiled to the EVM bytecode. What is the EVM programming language? The advantage to work on the Ethereum virtual machine allows to make transitions from one state to the other. For example, transactions get executed and gathered into a new block, the new state gets stored and this new block describes the next state. Here Merkle Patricia trees are used to optimize the verification process of

transactions on the EVM [Bisade, 2018].

The mapping between account addresses and the equivalent state of the account has to get stored within the Ethereum platform. That way the data doesn't need to be stored in a giant block header. The Merkle Tree structure is thus used to make the whole system scalable. In general, a Merkle is following the idea to have a big amount of data which are hashed together. They get split into smaller chunk buckets, this gets repeated until there is just a root hash left [Butin, 2015].

As stated before, this whole process is done to make it possible to perform Merkle proofs. They are much quicker and therefore better scalable. The proof consists of all the hashes going from the last one up to the root hash, if this gets verified by other persons, at least for the verified branch. The idea is, that within the database the data is stored in a Merkle tree and the main root is publicly known and proofed. If so, a person who wants to get the knowledge about data at a specific position this can be done within less time. Because just small amounts of data have to get searched [Buterin, 2015]. This capability to store information in Merkle trees is very useful to improve the scalability and it is called light nodes. So, in the blockchain system exists two types of nodes, full and light ones. Whereas the complete one does synchronize the whole chain and is also executing everyone. Even though a node does keep every transaction, it is not necessary to store all the data. One would just get the header of every transaction done since the first block was mined. This system works, because the data is stored in Merkle trees and can get verified top down. In conclusion, the advantages of using such a tree is that the hashes in one tree always dependent on the other hash buckets in the rest of the tree. Therefore, if someone wants to manipulate something, it is hard to accomplish. As the root hash always keeps information about actual the state, the receipts and the completed transactions, nodes can validate little parts of the chain [Kasireddy 2017].

After giving insights into the processing and storing process in the Ethereum World, the applications on the platform are another important topic. The Smart contract can be used to run code on the Ethereum Virtual machine in a decentralized way. It is comparable to real world contracts. A collection of conditions and actions, if a condition occurs, the associated actions are getting executed. As a result many real-world applications can be applied to smart contracts. The execution happens on the blockchain and is immutable ever after. One example often used, is a rent contract in the real world. In the real world one would pay rent on a defined date and get entrance to a place for the paid time. A contract represented on Ethereum would look the same, the contracts gets written in Solidity. Then the network executes the contract on the defined date and if the renter does pay as expected, means fulfill the condition he gets entrance to his place. If not, the contracts would simply not open the door anymore.

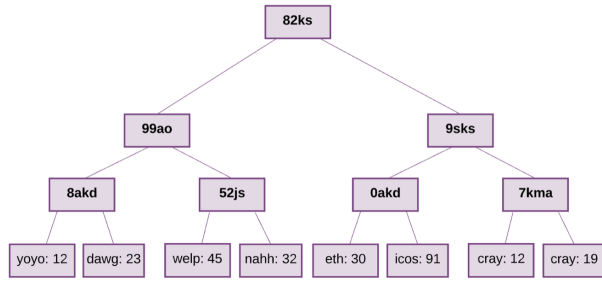
So, every component of a contract gets represented the enforcement, management and the payment component [Dannen, 2017]. Back to the example, the person who is renting the property would not get access to his door anymore, if

the rent was paid a day later. Here the first problem with smart contracts occurs, the rules are really strict, without any room for exceptional circumstances. Once the code is executed on the network it cannot be edited or corrected anymore. Everything gets executed as written in the code. If one does think about complex contracts this can be a problem. If something gets wrong or if there are paragraphs in it which are not covered by law, with a real contract the law is backing these contracts and one could take out paragraphs after a legal process. This is not the case in Ethereum, because to edit such a contract the entire network would have to get convinced, that something has to get changed. Which is almost impossible [Dannen, 2017].


Ethereum also provides a currency to run and incentivize the previously described system. In Ethereum the currency is called Ether. To execute and deploy contracts on the platform and to fulfill the contracts a fee has to get paid. This will be done in Gas. The concept can be subdivided into gas price and gas. The gas-price represents the amount of Ether that one is willing to pay for every unit of gas. The subunit of Ether is Wei and is representing the smallest unit. The Gas serves as a fee that is required to execute something in the network. As shown in the figure, every operation spends the defined amount of gas. In case the gas is not enough to fulfill the transactions, the sender runs out of gas and the End transaction does not get reached.

If someone wants to execute something in the network, one has to decide about the gas limit and gas price. This has an influence on the amount of computations that one transaction can do. The miner can collect transactions from the transaction Pool, after they pick them their respective code gets executed on the Virtual Machine and the Merkle proof gets created. The last step is to confirm the transaction with the proof of work algorithm. The reward depends on the selected amount of gas for the transaction. That also means, that some transaction might not get mined in case the fee is chosen to low. Another part the gas is used for, is the storing fee, the amount depends on the size of the storage used [Preethi, 2017]. On the one hand this was done, so that people will write optimized and efficient code and won't waste the Ethereum network computing power on unnecessary tasks [Preethi, 2017], [Dannen, 2017]. On the other hand, this is a way to incentivize the mining process of transaction. This helps to verify the transaction and should be described a little more into detail. This process verifies the transaction and is carried out to secure against counterfeiting. It is also applying the proof-of-work and uses the Ethereum network, respectively the nodes and creates new blocks with validated transactions.

The technology also has downsides, one is definitively the scalability. Every node in the system is storing the full data and executing every transaction. To store the state is a matter of security, also the high degree of decentralization. As more traffic gets on the system and more applications get written the data that have to be stored and verified for in the chain, get more and more. Specially on the Ethereum platform where there are more than just currency transaction



(a) Merkle Tree Structure



(b) Transactions

other than the Bitcoin. To be comparable to common systems, the Ethereum blockchain has to be able to keep on to the security and decentralization but also work on the scalability. So far, the community is working on different alternatives, such as *sharding* and *plasma chains* [Gadaleta 2018].

Vitalik Buterin describes Sharding, combined with Plasma as the most promising solution to make the platform more scalable [Giese 2018]. Here Sharding describes a partition of the main net in smaller shards. So far, the blockchain can support 15 transactions per second, with the new solution this could work for 100 transactions afterwards. Combined with the plasma solution this could scale up to a million transactions per second. The plasma solution represents a second-layer solution could bundle transactions and handle therefore more at ones [Giese 2018]. The main issue always is, that if one wants to higher the scalability, the verifications has to be reduced or less difficult. In this triangle one always reducing something, what means the points which have to be manipulated in the network are getting reduced. Next problem is the high computational power that is needed to verify or mine the transactions, here a lot of energy is needed that the system can work. So far, there is no solution for that.

But with the change to the proof-of-stake this should change. Besides the scalability and the energy consumption Ethereum has to deal with the high volatility in the crypto market. One can imagine, that in case of Ethereum the volatility has an influence on the usage of the platform. As described before the currency, called Ether is meant to be the Gas (fee) to pay the system for executing applications. If the price went up and down and is opposed by speculation, the original purpose might not get used anymore [Chandersekhar 2018]. Chandersekhar looked into the data of Smart Contracts and found out, that 94 % of the smart contracts are called less than 10 times and just 5 % are getting called between 10 and 100 times.

### III. ETHEREUM AND THE BLOCKCHAIN

Bitcoin which was 2008 introduced, is the more mature platform, compared to Ethereum. We want to give a short overview about both technologies and compare them. As the original paper is analyzing and describing the Bitcoin world. On the one hand we want to give some insights in the most important technologies on the blockchain market and also

might gather some information about possible problems with the outcome.

Lets start with the data, by now the Bitcoin blockchain has the size of 197 gigabytes in the beginning of January its growing constantly, since it was introduced [Portal, 2019b]. It takes on average 9,6 minutes to confirm a transaction at the end of February 2019 [Portal, 2019b]. The Bitcoin index value for the same period was 3.799,68 US Dollar, were the number of wallets grew over the last years constantly and reached 32 million wallet users at the end of 2018.




Fig. 4: Bitcoin Price Index- [Portal, 2019b]

The market capitalization had a peak in Quarter four 2017 and it has since decreased, to a value of 66,18 million US Dollar in the end of 2018. Compared the Ethereum already grew to the size of 181.65 GB, this represents the full blockchain size of all blocks. The average mining time is less than the Bitcoin, with 13.5 s per block [BitInfoCharts, 2019b]. The number of wallets exceeded 50 million in the end of 2018 [BitInfoCharts, 2019a].

Even if the wallets are increasing, the number of active addresses, referencing to BitInfoCharts, there are less active Ethereum addresses in 2018. The price of Ethereum is in the End of February 2019 amounted to 134,82 US Dollar [Portal, 2019a]. Comparable to the Bitcoin, the peak was in 2017 and has been decreasing since then. Similar to the peak in price to




Fig. 5: Ethereum Price Development - [Portal, 2019a]

most active period was in the end of 2017, here more than 1.13 million addresses were registered [Partz, 2018]. In the End of 2018 there were only 328.400 active addresses recorded.

The data can give some insights, but just in relation to the idea of both technologies they get more relevance. The most important difference is, that the idea behind the bitcoin was to create an independent currency. That is not easy to manipulate and does not need an intermediary to be trusted. The community should be the verification, that no one would manipulate the blockchain. The technology was described in the former chapter. In contrast the Ethereum was developed afterwards. The person who wrote the white paper, Vitalik Butin, wanted to create more than a currency. [Buterik, 2013] He got inspired by the technology behind the bitcoin and wanted to go beyond. That why he created a platform usable to write applications on, backed with the same technology. As described in the Ethereum Theory chapter, along the way he also was able to improve some downsides about the bitcoin, introducing the merkle Patricia tree. Another advantage is that he made the platform able to work with a turin-complete language. In the End all that was necessary, as the Ethereum Blockchain has to deal with way more complexity when it wants to serve as a platform for applications.

#### IV. LITERATURE REVIEW

This paper is inspired by the study of the Bitcoin transactions network performed by [Lischke and Fabian, 2016]. There, the authors examine the data for the first 4 years of the cryptocurrency with a goal of getting insights about the underlying economy and its evolution. They also apply a number of network metrics in order to identify major nodes, clusters, and verify whether the small world phenomenon exists for the Bitcoin network.

The main difference of our research is that it targets Ethereum cryptocurrency. We replicate some of the approaches suggested in [Lischke and Fabian, 2016] and adopt the others or come up with new approaches due to the differences in the

cryptocurrencies' design and supporting infrastructure. More details on this are given in the Methodology section.

To make our analysis grounded on the previous Ethereum research, benefit from findings and approaches described by other authors, have an opportunity to compare our findings to previously obtained results and potentially identify blank spots in the previous studies we conduct a literature research.

As Ethereum was introduced only in 2013 the literature landscape is relatively scarce. Most information is not peer-reviewed and published on various blog-platforms, forums and wiki-resources. This does not necessarily imply pure quality. However, we decided to use those publications only as a source of background information. Instead, the focus was made on publications that can be found using Google Scholar web search engine that indexes the full text scholarly literature.

We identified 361 articles that are published in the period from January 2013 till March 2019 and have the key word "Ethereum" in their title. For comparison, the search for articles published in the same time range but with keywords "cryptocurrencies", "bitcoin" and "blockchain" returns 618, 3670 and 6540 results respectively. The dynamics in the number of publications that target Ethereum is presented on the Figure 6a. Results for both broad and narrow (in combination with other keywords) searches are presented. There is a considerable growth over the last years. Nevertheless, only 11 % of articles is aimed at the Ethereum blockchain analysis. Thus, using a combination of the keyword "Ethereum" with keywords "analysis" and "network" we identified 23 and 16 publications respectively (see Figure 6b). Using other keywords such as "description", "graph", etc. provided only 1 additional article.

Our literature research is based on (but not limited to) these articles. At the later stage we excluded part of them from further review as they appeared to be out of immediate scope - were dedicated to analysis of price fluctuations, private implementations of Ethereum blockchain, new utility tokens or blockchain applications, social infrastructure around the Ethereum blockchain (like meetups or forums for blockchain discussions), etc. More publications were identified while studying the papers and by screening the search results for a broader keywords such as "Ethereum" or "cryptocurrencies".

Most of the literature corpus relevant to the topic of Ethereum blockchain analysis can be grouped in three broad categories: (i) data collection, (ii) descriptive analysis, (iii) network metrics and graph analysis, (iv) smart contracts.

*Data collection.* Certain part of the literature corpus is dedicated to the ways how Ethereum blockchain data can be gathered and aggregated for the sake of the future analysis. For instance, [Li et al., 2017] develop a query layer that allows collecting various data about blocks and transactions, as well as calculating balances of accounts. [Pratama and Mutijarsa, 2018] come up with similar query functionalities. They facilitate access to blockchain data based on multiple search parameters (retrieval query), provide simple statistical analysis (aggregate query) and sort data according to its various components (ranking query). [Pérez Jiménez, 2018]




Fig. 6: Results of Google Scholar search for articles.

analyzes different node client software and APIs that can be used to access blockchain data. He develops a tool that serves as an interface between the blockchain data and data analysis packages in Python.

*Descriptive Analysis.* A number of publications is focused on performing the descriptive analysis of the Ethereum blockchain. Thus, [Payette et al., 2017] research Ethereum address space. Using the k-means clustering algorithm and other techniques they identify 4 distinct groups of addresses, evaluate them quantitatively and qualitatively. They show that clusters differ significantly in terms of the volume and amount of ongoing and outgoing transactions. Authors also suggest that analysis of the amount of gas spent is crucial in "determining whether Ethereum is being utilized as a distributed supercomputer as intended by its creators or a vehicle for financial speculation". [Anoica and Levard, 2018] perform quantitative description of activity on the Ethereum blockchain. They estimate distribution of internal activity among users, identify most active addresses. According to the authors only 3 % of nodes have participated in more than 10 transactions, while smart contracts account for only 9,5 % of value transfers. [Kim et al., 2018] also evaluate activity of the Ethereum nodes. They find out that less than a half of all nodes contribute to the Ethereum Mainnet. A comparison with other P2P networks such as BitTorrent or Bitcoin reveals that Ethereum differs in both network size and geographical distribution. Ethereum is accessed as a small and immature P2P network that would benefit from adoption of known best practices. Authors also identify network inefficiencies such as outdated clients susceptible to patched vulnerabilities and hard fork incompatibility.

Descriptive analysis is done with a second goal of security assessment in several papers. For instance, [Kiffer et al., 2017] research the Ethereum/Ethereum Classic hard-fork that happened in 2016. Based on comparison of various descriptives, such as hash rate, difficulty per block, number of transactions, number of blocks won by the major mining pools, the number

of rebroadcast transactions, etc. they show the impact on the two networks and their mining pools. They suggest that the fork lead to unintentional incentives and security vulnerabilities. For instance, there is a security vulnerability wherein attackers can rebroadcast transactions from one network into the other. Ethereum Classic experienced a loss of almost 90 % of the nodes in its network immediately after the fork which also exposed it to various types of attacks.

Sometimes analysis is accompanied by improvement or application suggestions. So, [Bentke, 2017] researches pending transactions and optimal acceptance criteria for them. Author shows that pending transactions are related to the problem of double spending attacks and issues with the network topology. He suggests workarounds for these problems that are applicable in certain circumstances. [Dennis and Disso, 2019] analyze historic data about Ethereum and Bitcoin networks. Based on this analysis they model the growth of these networks in the future and propose ways that can help to reduce the networks' size and increase their scalability. [Signer, 2018] develops a tool to analyze and visualize gas costs and consumption. Author suggests that his tool can be used to explore gas consumption before deployment of smart contracts, is useful for a more gas-efficient and secure development of them.

*Network metrics and Graph analysis.* Many articles are based on the network theory and apply specialized visualization tools in their research. Thus, [Chen et al., 2018] perform graph analysis in order to evaluate Ethereum major activities: money transfer, smart contract creation, and smart contract invocation. For each type of activities they estimate network metrics such as degree distribution, clustering, degree correlation, node importance, etc. Based on a set of quantitative estimations authors conclude that (i) Ethereum is used mainly for transactions, smart contracts are not adopted actively, (ii) many of the users don't use Ethereum frequently, (iii) a considerable number of smart contracts is created by a small number of developers, (iv) financial hubs, such as exchanges dominate Ethereum. Besides, they showcase



how graph analysis can be used (i) to detect abnormal and potentially malicious contract creation, (ii) to identify smart contracts and accounts controlled by the malicious developer. [Chan and Olmsted, 2017] suggest a model for collection and graph analysis of the Ethereum transactions through a graph database management system Neo4j. Authors collect from Etherscan block explorer a set of addresses that are known to be associated with hacks on currency exchanges Gatecoin and Coindash. Using the model they try to de-anonymize these addresses and find out that hacked money ended up at a set of cryptocurrency exchanges (Changelly, ShapeShift, Bitfinex, Bittrex, and Poloniex). Authors suggest that linking Ethereum and Bitcoin transaction graphs in order to follow the money further and potentially reveal identities of the hackers.

Network theory based articles usually evaluate power law distributions in Ethereum network. For example, [Somin et al., 2018b] and [Somin et al., 2018a] study the dynamics of the social signals of Ethereum network, provide insights about the ecosystem and the forces acting within it. Authors demonstrate that the network displays strong power-law properties. According to them, this is the case for degree distribution of nodes, but as well for distribution of token popularity, in terms of buyers and sellers amount. [Anoaica and Levard, 2018] conduct similar research and conclude that the network is robust to isolated attacks as power law distribution of nodes degree lies in the range that is observed for most real networks. [Zwang et al., 2018] use network theory to detect bot activity on Ethereum Blockchain. They show that time difference between consecutive Ethereum transactions also follow the power law distribution. Consequently, deviation from this distribution are helpful to identify non-human activity.

*Smart contracts.* Given the role that smart contracts play in Ethereum blockchain we also consider publications on this topic. For instance, [Bartoletti et al., 2017] research practical applications of smart contracts. They identify different categories of smart contracts such as "financial", "game", "wallet", "library", etc. and quantify their usage. Authors also come up with a classification of the most common design patterns and relate them to application domain categories. [Bistarelli et al., 2019] analyze the code of the verified smart-contracts and compilers. They identify opcodes - functionalities that are of special practical importance - and perform a statistical research of them. Results are suggested to be used for optimization of the smart-contracts development. [Bennink et al., 2018] study how interaction between different Ethereum blockchains and tokens can be performed. They analyze atomic swaps - smart contracts that allow executing the cross-chain exchange of funds or single-chain exchange of different tokens in a decentralized fashion. Authors design single-chain token swap contract as a proof of concept, suggest theoretical improvements for other forms of atomic swaps.

A number of authors suggest specialized tools for analysis of the smart contracts. For example, [Albert et al., 2018] come up with a framework for decompilation of the EVM bytecode. The framework can be useful when the source code is not available. This is the case in many situations, for instance gas

consumption is specified at the level of EVM instructions. The framework allows performing analysis of smart contracts, for example in order to infer their loop bounds. [Grishchenko et al., 2018a] develop a tool for bytecode analysis that captures important security properties of smart contracts, such as single entrance or independence of the transaction environment, and help to identify when they are violated.

The topic of security concerns is dominant in articles dedicated to smart contracts. Thus, [Grishchenko et al., 2018c] do semantic analysis of Ethereum smart contracts. They define security properties of smart contracts, such as call integrity, atomicity, independence of transaction environment, independence from miner controlled parameters, etc. Authors identify mistakes and impressions in existing semantics and verification tools for Ethereum smart contracts, suggest their improvement. This research is further extended in [Grishchenko et al., 2018b], where authors overview the state-of-the-art in smart contract verification. They compare different verification tools for automated bug-finding, verification of generic properties, semi-automated proofs for contract specific properties, dynamic monitoring for predefined security properties, etc. [Tikhomirov et al., 2018] come up with a classification of code issues in Solidity language used for creation of smart-contracts. They also suggest a statistic tool that can help to detect these issues. Based on the large corpus of smart contracts authors find out which of them are the most common and severe. [Bartoletti et al., 2017] suggest ways to identify Ponzi schemes on Ethereum blockchain. Authors analyse their development and impact on participants and Ethereum blockchain from various viewpoints. According to the research 10 % of the studied 1384 contracts with verified source code on etherscan.io are Ponzi schemes. Nevertheless, the impact of these smart contracts is limited, as they cover only 0.05 % of the transactions on the Ethereum blockchain.

A meta analysis of the articles shows that all 4 major categories are covered in a relatively equal way (see Figure 6a). Data collection about the blockchain has less attention of specialized articles. This seems to be reasonable given that data collection is just the first step of the analysis and other articles usually also have some minor part dedicated to this process. Most of the attention is given to smart contracts, which can be explained by the importance of this Ethereum blockchain feature.

In our research we will focus on descriptive and network metrics that are slightly less present in the current research landscape. In conformity with the best practices set by reviewed papers, we will aspire to perform our analysis having in mind potential security and privacy issues that can be identified or inferred in the process.

To conclude the literature review section we should add that a number of other articles that are not directly related to the analysis subject deserve attention of those interested in future research on it. Thus, we have considered generic information about Ethereum: including its white and yellow papers (see [Buterin et al., 2014] and [Wood, 2014]) that describe technical specifications and design features of the blockchain; compar-

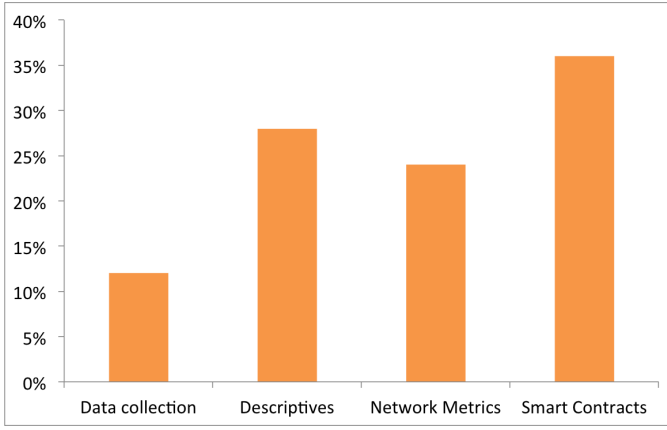


Fig. 7: Literature Landscape.

ative studies of cryptocurrencies (see for instance [Maesa, 2018], [Rudlang, 2017], [Sapuric et al., 2017], [Anderson et al., 2016]) that allow to get better understanding of the differences between Ethereum and Bitcoin blockchains and provide necessary background for analysis and reasoning.

To complete the picture we also considered research publications on network analysis of Bitcoin blockchain that are cited in [Lischke and Fabian, 2016] and have provided us with some additional insights and inspiration: [Reid and Harrigan, 2013], [Baumann et al., 2014], [Drainville, 2012], [Ober et al., 2013], [Meiklejohn et al., 2013], [Spagnuolo et al., 2014], [Androulaki et al., 2013], [Kaminsky, 2011], [Ortega, 2013].

## V. METHODOLOGY

### A. Data Preparation

In order to access the data in any form of blockchain technology one must first be an owner of a full node of the respective network. This means that the whole blockchain data with all its records is stored on the researcher's device. An overview of the information stored in the Ethereum blockchain is given in figures 8a and 8b. However for the most part, data in this form is not stored in a manner that is easy to query. For this reason we opted to become owners of a full node of the Ethereum blockchain, to subsequently scrape the blockchain for the desired information and to store the results in a relational database, which in turn is easy to query.

The data-scraper is written in the Python language using the `web3.py` module [Rossum, 1995], [Piper and Caver, 2018]. Through `web3.py`'s interface it is possible to interact with the Ethereum blockchain. In particular the `web3.py` module enables us to extract information for every block and each transaction that is recorded in the blockchain.

In particular, for each block inside a predetermined range we query its contents and obtain the data that is interesting for the further analysis. On the block level we extract the hash of each block as a unique identifier and the time each block was inserted into the blockchain. Note that this is not necessarily the time that the transaction was being commissioned. Finally we query the total gas that was used for each block as well

as the gas limit for each block. Along with that information we retrieve a list of transactions that are contained within that block.

Making use of the list of transactions within a block we extract the wallet addresses of both sender and receiver, as well as the total value, of the transaction. Additionally we retrieve information on the gas price for the transaction. Figure 9 summarizes the relations within that database and Algorithm 1 gives an outline of the datascraper in pseudocode.

---

#### Algorithm 1 Data Scraper: Overview

---

```

Check progress file if the ETL can be resumed from a block.
if yes then
    current block  $c \leftarrow$  content of progress file
else
     $c \leftarrow 1$ 
end if
 $s \leftarrow 2000000$ 
while  $c \leq s$  do
    Overwrite progress file with  $c$ 
     $content_{Block} \leftarrow getBlock(c)$ 
     $row_{Block} \leftarrow extractInfo(content_{Block})$ 
    write( $row_{Blocks}$ , "Block")
    for each  $transaction$  in  $content_{Block}$  do
         $row_{Transaction} \leftarrow extractInfo(transaction)$ 
        write( $row_{Transaction}$ , "Transactions")
    end for
end while

```

---

### B. Descriptive Analysis

We collect the basic descriptive metrics that allow getting the initial insights about the network structure and interactions between nodes. Part of these metrics describe distributions as well as average, median, minimum and maximum values for:

- value of a transaction,
- amount of gas used in a transaction,
- fee paid for a transaction,
- number of transactions per block,
- value of transactions per block, etc.

Another part follows the changes that happen to the network across time. Those are:

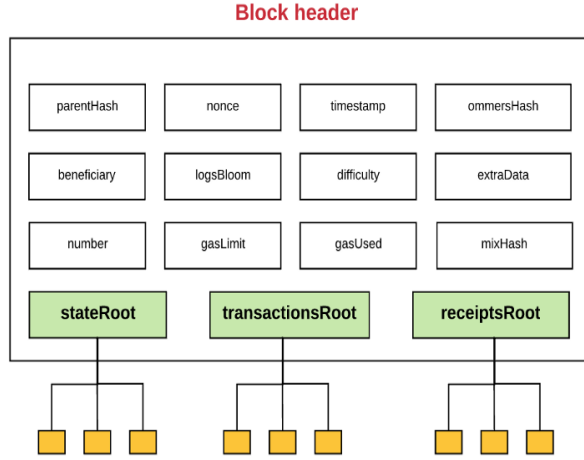
- number of transactions over time,
- volume of transactions over time,
- average fees paid per transaction over time, etc.

We associate transactions with individual addresses. This helps to identify:

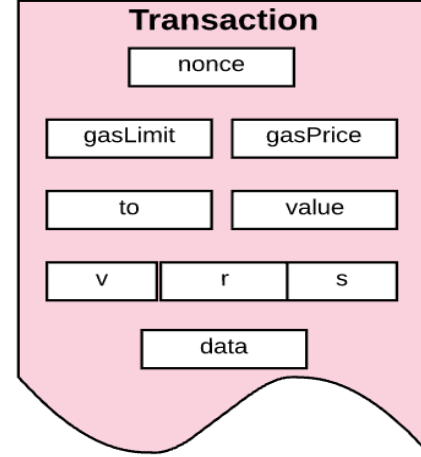
- major senders and receivers of transactions (both in terms of the number of transactions and monetary value exchanges in these transactions),
- net balances of accounts in a studied period of time,
- distribution of the net balances across the network,
- percentage of wealth accumulated and transacted by major nodes, etc.

In order to understand our findings in more detail we try to find out identities of the major nodes using data of





(a) Block headers [Preethi, 2017]



(b) Transaction contents [Preethi, 2017]

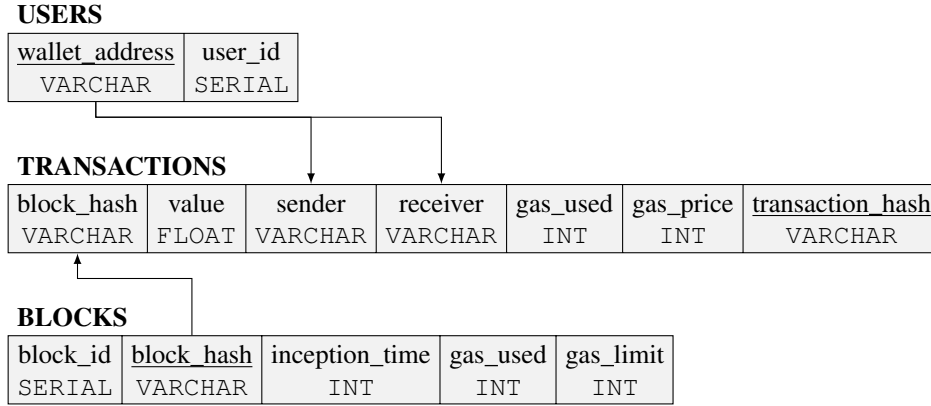


Fig. 9: Database Model: the primary keys of each table are underlined, foreign keys are represented by arrows, the type of each column is capitalized.

the Ethereum scanners such as Etherscan.io, Etherchain.org, Blockchair.com, Bloxy.info, etc.

When relevant (for example in case of transactions-outliers, atypical behavior of nodes, discrepancies in data, etc.) we also consider Ethereum forums, posts and newsfeed in order to get additional information about events and actors that may explain or discard certain findings.

### C. Graph Analysis

We can characterize a network by a number of metrics. The degree of a node in the network is the number of incoming or outgoing edges. We define a *triplet*  $i, v, j$  as an ordered set of three nodes, where  $v$  is the focal point and the undirected edges  $\langle i, v \rangle$  and  $\langle v, j \rangle$  form the neighboring edges. A triplet is considered closed, when there are exactly three connections for the three nodes. Three closed triplets form a *triangle* [Meghanathan, 2018]. Additionally, we can define multiple measures to investigate a node's role inside the network.

1) *Degree Centrality*: The most common centrality measure focuses on counting the number of nodes in it's immediate

vicinity. It measures the centrality by counting the number of direct connections to the node. For a network with  $g$  nodes, the degree centrality of node  $n$  is defined as

$$C_D(n) = \frac{k(n)}{g-1}, \quad (1)$$

where  $k(n)$  denotes the degree of node  $n$ . The Degree Centrality is normalized by  $g-1$  since this is the highest possible degree, a node that is connected to every node in the network other than itself [Meghanathan, 2018].

2) *Closeness Centrality*: While the Degree Centrality only takes the adjacent nodes into account, it may be possible that the most central in this sense, i.e. the node with the highest degree, is in fact not close to other nodes in the network. This shortcoming is overcome by the *Closeness Centrality*. It considers the sum of geodesic distances, i.e. the number of edges in the shortest path between two nodes. The Closeness Centrality of a network with  $g$  nodes is defined as

$$C_C(n) = \frac{g-1}{\sum_{i=1}^g d(n, i)}, \quad (2)$$

where  $d(i, j)$  stands for the geodesic distance between node  $i$  and  $j$  [Meghanathan, 2018].

3) *Betweenness Centrality*: So far, all the measures focus heavily on the direct influence of the node in question on others. However, two nodes which are not adjacent can also influence each other indirectly through other nodes in the network. The *Betweenness Centrality* highlights the nodes that lie most frequently in the connecting path between two nodes. For a network with  $g$  nodes the Betweenness Centrality for node  $n$  is defined by the sum over all geodesic distances that pass node  $n$ , normalized by the sum of all shortest paths of all pairs of nodes:

$$C_b(n) = \frac{\sum_{j < k} g_{jk}(n)}{g_{jk}}.$$

Here,  $g_{jk}$  denotes the number of shortest paths between two nodes of the network and  $g_{jk}(n)$  marks the number of shortest paths that pass through node  $n$ . Since the maximum number of such paths is

$$\frac{(g-1)(g-2)}{2}.$$

One can thus further normalize:

$$C_B(n) = \frac{C_b(n)}{\frac{(g-1)(g-2)}{2}}. \quad (3)$$

4) *Clustering Coefficient*: In order to account for local structures in graph theory, one uses the *Clustering Coefficient*. There is evidence that suggests that the ties between nodes in most real-world networks tend to create tightly knit groups with a relatively high density of connections [Watts and H. Strogatz, 1998].

There exist two versions of this measure, the *global* and the *local* Clustering Coefficient. The global Clustering Coefficient  $C_G$  is the number of triangles over the total number of open or closed triplets in the graph. This measure can be applied to both directed and undirected networks.

The local Clustering Coefficient for a node  $n$  is given by the proportion of links between all adjacent vertices divided by the number of links that could possibly occur between them. For a directed graph the in- and outgoing edges have a different meaning and therefore the maximum number of immediately adjacent neighbors of  $n$  is  $k(k-1)$ , where  $k$  is the degree of node  $n$ . The neighborhood of a node  $n$ , i.e. its immediately connected neighbors, is defined as

$$N = \{n : e_{ij} \in E \cap e_{ji} \in E\},$$

where  $E$  is the set of edges.

Thus the local clustering coefficient of node  $n$  is given by:

$$C_L(n) = \frac{\text{number of closed triplets centered around } n}{\text{number of triplets centered around } n} = \frac{|\{e_{jk} : v_j, v_k \in N \cap e_{jk} \in E\}|}{k(k-1)}. \quad (4)$$

In an undirected graph the meaning of in- and outgoing edges is considered identical. Therefore if a node  $n$  has degree  $k$ , there could exist a total of  $\frac{k(k-1)}{2}$  edges within  $n$ 's neighborhood. Thus the local Clustering Coefficient of a undirected network can be defined as

$$C_L(n) = \frac{2 \cdot |\{e_{jk} : v_j, v_k \in N \cap e_{jk} \in E\}|}{k(k-1)}. \quad (5)$$

This means that nodes with a high Clustering Coefficient play a large role when relaying information through the network [Knorr, 2019].

5) *Power Law and the Small World Phenomenon*: The small world phenomenon describes that the size of a network does not have any effect on the size of ties among its nodes. This means that the path lengths are characteristically small just like random graphs [Watts and H. Strogatz, 1998].

In a random graph the degree of a node  $n$  follows a power law distribution, i.e. we can say:

$$P(K(n) = k) \propto k^{-\alpha}, \quad (6)$$

where  $K(n)$  is the random variable that describes the degree of node  $n$  and  $\alpha$  is a constant whose values range between 1.6 and 3.0 [Newman, 2003]. This distribution suggests that most nodes will have a low degree, while a small number of nodes will display a large number of connections [Alstott et al., 2014].

## VI. FINDINGS

### A. Data Preparation

We encountered many difficulties during our attempts at collecting and refining the data. At first, we attempted to deploy the data-scraper on a RaspberryPi 3 Model B in conjecture with an external HDD hard drive. We decided on this option, since it did not require any of us to commit indispensable computing resources and at that time it seemed to be the most convenient option.

That attempt was, in hindsight, doomed to fail due to the hardware's limitations. The problem did not lie in the hardware intensity of the data-scraper in itself, but rather in the inordinate amounts of memory it takes for `go-ethereum` to synchronize with the Ethereum blockchain. The interface to the ethereum blockchain, `go-ethereum`, is written in a manner that is primarily optimized for mining ether and in order to facilitate the commissioning of transactions and smart-contracts. The way the program achieves this is by consuming as much memory as there is available on the system, or at least as much memory as the user is willing to allocate to it before starting the program. However, whenever the program runs out of memory or whenever it attempts to

exceed its predetermined limit the process shuts down and kills itself. This does not result in a system crash, but it does inhibit the workings of the Python libraries that are used inside the data-scraper.

This design of the `go-ethereum` program frequently led to the `RaspberryPi`'s 1 GB of memory being filled after only a short period of time, usually between thirty minutes and three hours. Note, that the data-scraper's construction did anticipate interruptions and it is possible to resume scraping the Ethereum blockchain from the current block without the loss of any data and without the need to redundantly query more than the one block at which the disconnection happened. This behaviour meant that the `RaspberryPi` was idling for most of the time.

Unfortunate as this was, it was not the only problem. The use of an external hard disk drive for storing the database meant that a considerable amount of the time during which the `RaspberryPi` was productive was spent on input/output-operations while it was committing the scraped information into the database. In order to alleviate the bottleneck that was presented by the I/O-operations it would have been favorable to use an external solid state drive for that task. However, we were too frugal to obtain one just for the singular purpose of this seminar paper.

This means that after running the data-scraper for an entire month, along with taking all the care that is reasonable, in order to keep the scraper afloat, we were able to extract 120.000 transactions from 15.000 blocks which tantamounts to merely three days of data.

We solved all these problems by taking up the offer *Microsoft Azure* provides for students. Using this free initial student credit on the *Microsoft Azure* servers it was possible to set up a virtual machine with 32 GiB of memory along with access to an additional solid state drive. The fact that it is possible to rent out a Linux machine on the *Microsoft Azure* servers made it possible to migrate the scripts for the data-scraper without the need for adjustments.

TABLE I: Summary statistics

	value	gas used	gas price
mean	48.12	127511.76	3.01e+10
std	1495.49	239571.52	1.34e+13
min	0.00	21000.00	0.00
25%	0.06	40000.00	2.00e+10
50%	0.89	90000.00	2.00e+10
75%	1.11	90000.00	2.50e+10
max	1000000.00	4712388.00	3.62e+16

The initial student credit was enough to deploying the data-scraper on the *Microsoft Azure* servers for about a week. During this week we were able to extract 2 GB of data that consist of one million blocks with a total of 7.3 million transactions made from 415000 distinct wallets. Table I describes some rudimentary statistics.

Initially we also planned to scrape the IP-addresses of a wallet by making use of `web3.py`'s functionality which is provided for the `Bitcoin` blockchain. However, within

the scraped dataset we were not able to identify a single address. The reason might that the `web3.py` module does not provide that functionality for the Ethereum blockchain in its current version. Mind you that we do not claim that it is impossible to retrieve the IP-address of a wallet. One proposed heuristic to do so is to assume that the first node that propagates a transaction through the network is also the one that is initializing the transaction [Reid and Harrigan, 2013]. However, for us this would have meant to rewrite the data-scraper from Python into the Go language, which we considered to be out of this analysis' scope.

## B. Descriptive Analytics

Descriptive analysis of the data allows getting insights that can be useful at the time of the network analysis. By itself alone it can point at important information with regard to the network structure, interactions between the nodes, underlying business processes, design problems, potential privacy and security issues.

For instance, while the median value of a transaction is equal to 0.9 ether, the value of the biggest one is 1 million ether (138 million U.S. dollars as of 15.03.2019). Figure 10 shows the distribution of value sent in individual transactions. There is a significant number of outliers indicating at certain concentration of ether holdings.




Fig. 10: Distribution of Value sent per Transaction.

\*The number of transactions on the Y axis is limited to 50, so that the less frequent values remain visible.

The spikes in the amount of fees paid in transactions (see Figure 11) gives a hint at the underlying variation in businesses activity and Ethereum design issues. Cryptocurrencies are often promoted as a fast medium for micro-payments. However, fast transactions with low fees are not always the case. Limitations in the block size and the number of blocks in a period of time may result in an overflow of transactions pending to be included in a block. To speed up transactions people pay higher fees. At the same time, when the exchange rate of a cryptocurrency rises, fees tend to become lower.

Interestingly, there is a big outlier - once there was a fee of more than 600 ether (80 000 U.S. dollars as of 15.03.2019). Apparently this is due to a mistake made by a payer. Such




Fig. 11: Distribution of Fees paid per Transaction.

*\*The number of transactions on the Y axis is limited to 700, so that the less frequent fees remain visible.*

a mistake shows one of the disadvantages of the blockchain-based currencies where all transactions are irreversible.

Analysis of the distribution of gas used in transactions (see Figure 12) points at differences in their types. As gas measures how much "work" an action or a set of actions takes to perform, this potentially can be used to make a preliminary judgement about the prevalence of the smart contracts or about distribution of their complexity.

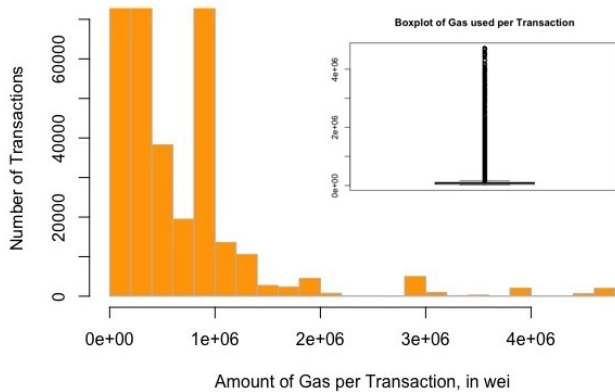


Fig. 12: Distribution of Gas used per Transaction.

*\*The number of transactions on the Y axis is limited to 70 000, so that the less frequent amount of gas remain visible.*

There are obviously two major clusters that probably separate regular payment transactions and smart contracts plus several groups of more complex contracts that however represent minority in the total volume of transactions. Both median and the 3rd quartile amount of gas is equal to 90 000 wei. In case the above hypothesis is correct, regular payment transactions will cover at least 75 % of the total number and are all represented on the Figure 12 by the 1st bar (cut in order to make other values visible).

Analysis of the transactions number over time makes the concern about unequal distribution of wealth in the network visible. As illustrated on the Figure 13 the number of transac-

tions on the most active days more than doubles the respective number on the days with the lowest activity. Such significant variation indicates at the presence of major nodes.




Fig. 13: Number of Transactions over Time.

This becomes even more visible if we consider the monetary volume of transactions over time (see Figure 14). Increases in volume of up to a hundred times from one day to another are an example of activity that would lead to a drastic changes in prices if these transactions are made not on chain but on currency exchanges.






Fig. 14: Volume of Transactions over Time, in ether.

In order to verify the hypothesis about the presence of major nodes we analyse transactions number and volume more in depth. This is done by associating transactions with individual accounts. We treat receivers and senders of transactions separately.

As shown on the Figure 15a top 5 major accounts have received 19,9 % of all transactions (1,45 out of 7,3 million transactions). Top 50 accounts cover slightly more than one third of the total transactions' number. When it comes to the monetary value, the situation is even more centralized. Top 5 accounts have received 31,8 % of the total transactions' value



(a) Top 50 Receivers.



(b) Top 50 Senders.

Fig. 15: Cumulative Percentage of Transactions (blue) and Volume (orange) covered by Major Receivers (a) and Senders (b).

(112 out of 352 million ether received in all transactions). Major 50 Receivers cover more than a half of the volume transacted within the network.

Figure 15b shows the data for major senders of transactions. Here, in terms of the transactions' number the situation is even more centralized. Top 5 accounts have made 36,4 % of all transactions. Top 50 senders cover 60,1 % of the entire transactions' number. At the same time in terms of volume the picture is different. Top 5 and Top 50 senders cover only 18,3 % and 35,0 % of the total transactions' value respectively.

Interestingly, there are significant discrepancies in the degree of centralization if the data about receivers and spenders is compared to each other. This is the case for centralization both in terms of transactions' number and volume. Analysis of these discrepancies may shade light on some of the roles that major nodes play in the network.

*Discrepancies in the Transactions' Volume between Receivers and Senders.* There is a significant difference between the amount of ether received and spent by the major accounts. While Top 50 Receivers have accumulated 172 million ether, Top 50 Senders have spent only 123 million. Besides, top receivers and senders should not necessarily be the same entities. Consequently there is a number of major nodes that spend less than their receive or don't spend at all. Figure 16 shows the net balances of the accounts within the studied period of time. While absolute majority of accounts have balance around 0, there are few that have accumulated a significant amount of wealth.

Those could be mining pools that receive new coins and hold them anticipating appreciation of ether. However the amount of Ethereum accumulated by such nodes seems to be too big to be explained this way. Despite a new block in Ethereum is mined every 10-20 seconds and remunerates miners with 5 ether, during the analyzed period of time this results in only approximately 5 million new ether.

Probably, these nodes are currency exchanges that accumulate wealth of their participants. They keep track of exchange transactions internally without recording them on-chain. The wealth is recorded as spent only when a participant makes a payment transaction.




Fig. 16: Distribution of Balances within Studied Period.

\*The number of wallets on the Y axis is limited to 10, so that the less frequent wallets remain visible.

If the nodes are actually currency exchanges, there is a deviation from the Ethereum narrative of "no trusted third party involvement". Given the level of centralization associated with the nodes, this could imply a major source of vulnerability. An attack on the exchange records would lead to losses for a significant number of the exchange participants.

In order to verify how relevant this risk actually is, we conducted further investigation of the accounts that received the biggest volumes of ether. Using data of the Ethereum scanners such as Etherscan.io, Etherchain.org, Blockchair.com, Bloxy.info, etc.<sup>1</sup> we found out affiliation of the Top 5 Receivers (in terms of volume).

Four out of five Top Receivers actually belong to the currency exchanges (see Table II). However, all these accounts in reality appeared to be not regular addresses but smart contracts. The remaining Top Receiver (the first one in the list) is also a smart contract. None of them is supposed to

<sup>1</sup>For details see:  
<https://etherscan.io/>,  
<https://www.etherchain.org/>,  
<https://blockchair.com/Ethereum/>,  
<https://bloxy.info/>.

TABLE II: Top 5 Major Receivers (Volume)

Wallet	Amount Received, mill. ether	Type	Affiliation
0xAA1A6e3e6EF20068f7F8d8C835d2D22fd5116444	66,0	Smart-Contract	ReplaySafeSplit
0xBFC39b6F805a9E40E77291afF27aeE3C96915BDD	13,2	Smart-Contract	Poloniex
0x209c4784AB1E8183Cf58cA33cb740efbF3FC18EF	12,6	Smart-Contract	Poloniex 2
0x7727E5113D1d161373623e5f49FD568B4F543a9E	11,1	Smart-Contract	Bitfinex 2
0xFa52274DD61E1643d2205169732f29114BC240b3	8,5	Smart-Contract	Kraken 4

TABLE III: Top 5 Major Senders (Number of TXs)

Wallet	Number of TXs send	Type	Affiliation
0xEA674fdDe714fd979de3EdF0F56AA9716B898ec8	733818	Address	Ethermine
0x2a65Aca4D5fC5B5C859090a6c34d164135398226	715017	Address	DwarfPool 1
0xD34DA389374CAAD1A048FBDc4569AAE33fD5a375	516932	Address	GenesisMining
0x52bc44d5378309EE2abF1539BF71dE1b7d7bE3b5	443544	Address	Nanopool
0x61C808D82A3Ac53231750daDc13c777b59310bD9	251566	Address	F2Pool 1

accumulate wealth and by design should transfer it further. This puts under question the validity or completeness of the transactions data we have collected.

Deeper research allowed restoring the faith in the dataset. Appeared that money received by identified major accounts was actually spent. Bit in another network.

The point is that on the 20th of July 2016 (shortly before the analyzed period) Ethereum experienced a hard fork. As a consequence of the DAO incident - a hack of a complicated smart-contract that resulted in a loss of about 12.7 million ether (worth around 150 million U.S. dollars at the time) - the community was split in two parts. Most of the network participants decided to make a fork chain where the stolen coins would be returned to their owners. However, some of the community decided to continue the old chain (Ethereum Classic or ETC), arguing that "the code is the law" and the blockchain data should be irreversible<sup>2</sup>.

This lead to a number of issues related to the interaction between chains. One of the most relevant of them was "replay attack". The mechanics of it is as follows: if there is a valid transaction on one chain, it can also be offered - replayed - on the other chain to duplicate the received amount. For instance, a user of the currency exchange can deposit and withdraw her money from this exchange on the old chain, and then use the 2nd transaction to withdraw money also on the new (forked) chain<sup>3</sup>. This concern was solved through inter-mediation of the smart contracts.

Thus, the post by Ethereum founder Vitalik Buterin states: "users who are interested in taking any actions with their ETC, including creating and participating in applications, converting to another asset, etc. are advised to use the splitter contract at address 0xAA1A6e3e6EF20068f7F8d8C835d2D22fd5116444 to move their ETC to a separate newly created account so as to avoid replay attacks; we also encourage the ETC community to consider adopting a secondary hard fork to change transaction

formats to make further replay attacks impossible<sup>4</sup>.

The address mentioned in the post is the biggest Receiver in our dataset<sup>5</sup>. It has transmitted all the received money to its senders, but on the other chain. The other 4 accounts, which are intermediate exchange wallets were apparently performing similar function of splitting wealth between different chains.

Thus, our concern about possible attack is confirmed only partially. None of the Top 5 Receivers were accumulating wealth and it all was automatically transmitted to other accounts. At the same time while exploring the "ReplaySafe-Split" smart-contract we have found that a number of users lost their money due to the mistakes made by them during the split<sup>6</sup>. The smart-contract itself appeared to be susceptible to a number of bugs<sup>7</sup> and was later substituted by another one. The DAO incident described in the context of our Receivers exploration also exemplifies how the risk of the attack on an account or contract that holds funds from different users can and has actually played out.

*Discrepancies in the Transactions' Number between Receivers and Senders.* As has been described by the Figure 15 the number of transactions received by major Receivers is significantly smaller than the number of transactions made by major Senders. This discrepancy can probably be partially explained by the activity of the mining pools. There, one account receives a mining reward and then sends it in small fractions to the many participants of the pool.

Alternative explanation can be related to various approaches of the network participants to gain more privacy. One of them is the so called "peeling-chain". The technique consists in dividing and sending the wealth of an account to multiple addresses again and again. The goal is to create an impression that several users are doing transactions instead of one.

<sup>4</sup>For details see [https://blog.Ethereum.org/2016/07/26/onward\\_from\\_the\\_hard\\_fork/](https://blog.Ethereum.org/2016/07/26/onward_from_the_hard_fork/).

<sup>5</sup>The original address in the post was substituted as the original smart-contract was changed.

<sup>6</sup>See for instance <https://medium.com/@chevdor/safer-version-of-the-replaysafesplit-smart-contract-a29c347e8a7>.

<sup>7</sup>For details see <https://etherscan.io/address/0xaa1a6e3e6ef20068f7f8d8c835d2d22fd5116444#code>.

<sup>2</sup>See for instance <https://www.cryptocompare.com/coins/guides/the-dao-the-hack-the-soft-fork-and-the-hard-fork/>.

<sup>3</sup>See for instance <https://vessenes.com/hard-fork-and-replay-concerns/>.



The process is illustrated on the Figure 17. Here, the received amount of 50 coins is splitted in several iterations to multiple accounts in an attempt to complicate tracking of a certain persons wealth.




Fig. 17: How the peeling-chain works. From [de Balthasar and Hernandez-Castro, 2017]

Further research of the Top 5 transaction Senders allowed verifying the hypotheses about the role of different nodes and their activities. Using data of the Ethereum scanners as before we disclosed identities of the major accounts. All of them appeared to be mining pools (see Table III. Among Top 5 Receivers of transactions there are no mining pools (the accounts are a digital token YoCoin, 2 Poloniex exchange wallets, currently closed by the U.S. government cryptocurrency trading platform BTC-e and already mentioned ReplaySafeSplit smart-contract). Thus the hypothesis about the activity of mining pools is confirmed. "Peeling-chain" practices are neither confirmed nor disproved.

### C. Graph Analysis

Looking at the rough outline of the observed period revealed that almost half of the wallets take part in one transaction during that time frame. Roughly 16 percent of the wallets appear two or three times respectively.

Since the vast majority of accounts only seem to appear once inside the whole dataset we concluded that these are not interesting from a graph analysis perspective. Therefore we opted to use a subset from that data of accounts, where we could verify that they are interacting with other accounts. This subset consists of roughly four and a half million transactions in which 27.000 distinct wallets were either the sender or the receiver. From this subset we opted to create a graph in which the wallets are represented by the graph's nodes and the transactions are symbolized by the graph's edges.

Even this subset proved to be too large to handle, which means for the further graph-analysis that we were forced to rely on a random subset. However, if this random sample is large enough it should still give us an insight into the structure of the Ethereum network and the graph formed by this subset should still provide interpretable results.

Figure 19 displays the graph built from a random subsample of 1000 transactions, where again the nodes represent the different wallets and the edges stand for an observed transaction between them. The size of the node stands for the transaction's value. We can clearly see that most of the accounts form connections that consist of only two wallets and it is highly likely that these are also the accounts which are only observed once in the entire dataset.


Even here it is possible to see that there are subclusters of highly connected nodes associated with a high transaction value located in the middle of the graph. This appears despite the small sample size, which has to mean that these nodes in the center of the graph are important actors inside the network whose position will most likely become more pronounced the larger we choose the sample size to be. From this first glimpse at the network we would expect to find two tendencies inside the dataset. First, we expect a large number of nodes which are associated with low centrality measures and clustering coefficients. Secondly we expect to find some collections of nodes that form subgroups and within these subgroups we should be able to locate some nodes that play a tangible role in their respective local neighborhoods. Both of those observations do fit into the notion of a random graph, where we would expect a large number of loosely connected nodes with a low degree and an exponentially lower quantity of nodes which have a higher degree and who play a more central role inside the network.

Note nonetheless, that figure 19 provides only a rough outline since it is only using one thousand out of the 7.3 million available transactions. We pick a comparatively small sample size for this graph in order to avoid overplotting, the further analysis however uses a larger subsample of 100000 transactions.


The results of the larger subsample are summarized in table IV and figure 18b. For the Degree Centrality we can observe a highly skewed distribution, where the mean lies at  $\bar{C}_D = 0.0001$ , whereas the first three quartiles remain firmly at 0. This indeed confirms our previous suspicion that most nodes inside this network are merely loosely connected, even after removing the most isolated and remote nodes of the initial dataset. This means that most wallets interact only with a very limited set of other wallets and that it is highly uncommon for a node to accept transactions from many different wallets.

When looking at the results for the graph's Betweenness-Centrality, we can see a similar pattern. The mean lies at  $\bar{C}_B = 0.0001$  and all three quartiles are, again, 0. This means that the Betweenness centrality's distribution is also heavily skewed, which indicates that only a few nodes act as a mediator for their neighbors. This should not come as a surprise, since most of the nodes have a very low degree to begin with.

The Closeness centrality between all nodes paints a different picture. Its distribution is still skewed, however it is not as extremely skewed as the measures that have been discussed so far. The mean lies at  $\bar{C}_C = 0.15$  with a standard deviation of  $s_{C_C} = 0.078$ . Note that this measure needs to be interpreted



(a) Distribution of total number of transactions



(b) Centrality Distribution

TABLE IV: Centrality Distribution

	Degree centrality	Betweenness centrality	Closeness centrality	Local clustering
mean	0.0001	0.0001	0.1540	0.0043
std	0.0012	0.0020	0.0782	0.0636
min	0.0000	0.0000	0.0000	0.0000
25%	0.0000	0.0000	0.1522	0.0000
50%	0.0000	0.0000	0.1876	0.0000
75%	0.0000	0.0000	0.2072	0.0000
max	0.1344	0.2149	0.2862	1.0000

with care. Most of the nodes are disconnected, and thus do not have a connecting path at all. If a node  $j$  cannot be reached from a different node  $k$  there exist two common propositions in the literature. Either  $d(j, k) = 0$  or  $d(j, k) = \infty$ . Most implementations rely on the former convention rather than the latter, since it avoids values that cannot be interpreted as a number. This is also the convention that is implemented in the `networkx` module for the Python programming language [Hagberg et al., 2008]. But this also means that the closeness centrality should be interpreted as a local measure since the distances in the normalizing factor in (2), which are calculated as

$$\sum_{i=1}^g d(n, i),$$

will contain numerous zero-valued summands. So actually, this measure only makes sense for nodes that are of a degree  $k$  that is larger than three, since nodes of a lesser degree are necessarily close to each other.

The Local clustering coefficient across the nodes is again highly skewed to the left. The mean lies at  $\bar{C}_L = 0.004$  at a standard deviation of  $s_{C_L} = 0.063$ . Again the quartiles are all planted at 0. Note however, that there exist some nodes whose Local clustering coefficient reaches as far as  $C_L(n) = 1$ . This means that there are some nodes who are fully connected within their local cluster, which means that all triangles within that clique are centered around that node  $n$ . This, in turn, means that there are some nodes which are perfectly connected inside their neighborhood.

So far all these findings indicate that the model of a random graph would be a good approximation of the real world data. In order to test this assumption in a graphical manner, we would like to compare the empirical density functions of the degrees against the theoretical powerlaw distribution. Remember that, in a random graph, the random variable  $K(n)$  that describes the degree of a node  $n$  follows a powerlaw distribution

$$P(K(n) = k) \propto k^{-\alpha},$$

where  $\alpha$  is some constant that has to be determined by Maximum-Likelihood methods [Meghanathan, 2018]. Figure 20 presents the results for the empirical probability density function (PDF) and the complementary cumulative distribution function (CCDF). The CCDF is primarily used in the domain of signal processing and measures the amount of time a signal spends above the mean power level of the measured signal. Equivalently, it assesses the probability that the signal power will be above the average power level [Udaigiriya and Sharma, 2015].

In order to reflect the difference in the roles across the network such as the receiver or the sender of a transaction, we split the analysis into three different parts. First, we look at the distribution of the directed graph that is focusing on the actors that initiate a transaction, i.e. the senders. Second, we look at the network formed by the receivers of a transaction and finally we look at the whole undirected graph that contains both receivers and senders. The results are displayed in figure 20. The PDF is shown by the blue line and the CCDF is represented by the red line. A high overlap between the




Fig. 19: Graph formed by a subsample of 1000 randomly drawn wallets

dashed and the solid line indicates a good fit of the empirical distribution compared to the theoretical. Thus the probability density function of a perfectly random network would lie on its respective dashed line.

In the leftmost part of figure 20 we can see for the senders that the tails of the empirical PDF are closely following the theoretical distribution of the random graph. However, in the central part we observe one large deviation from the distribution of a random network. In particular, this means

that - under a random graph - we would have expected more nodes that have a middling degree, i.e.  $k = 3$ ,  $k = 4$ , and so forth. Instead, we observe a large number of loosely connected nodes.

The same holds for the sender's CCDF. The lower tail of the distribution follows the random graph closely, whereas the central part is too low, and the upper tail is too high for a random graph.

Looking at the wallets at the receiving end of a transaction,




Fig. 20: Comparison with random graphs

in the middle part of figure 20, we can discern a similar pattern. Both, the line for the PDF, as well as the line for the CCDF begin by closely following the theoretical distribution and begin to deviate after passing the median. Again we observe a density at the lower tail that is very similar to what we would expect from a random graph.

However, the most clear indicator of that pattern is shown for the composite graph, that contains both senders and receivers, in the rightmost part of figure 20. Here we observe a very pronounced deviation from the course of both the theoretical PDF and CCDF after reaching the median. Again, this is a clear indicator that there are too many nodes with a comparatively high degree, for this to be a random network.

## VII. CONCLUSION

This work aims to transfer the approaches of the Bitcoin paper [Lischke and Fabian, 2016] to the Ethereum blockchain and get some insights into the Ethereum Economy. To do so, transaction data got scraped, displayed and analyzed.

Beforehand we looked into previous studies, to get a better understanding what was done so far and might find additional information about how to get good off-set data or IP addresses. As the Ethereum is still quite new, some papers were still published, which already looked into the Ethereum address spaces, security aspects and theoretical aspects of the Ethereum technology to find patterns within the network. None of them brought together transaction and network data to analyze patterns in the combined offset and network data over an extensive time period.

With that knowledge we worked our way through the same platforms, 'bitcoin.info' to enrich our data set and analyze it in the descriptive and graph analysis.

Even if we were not able to get data about the IP addresses, the descriptive analysis delivers information to get related to the network analysis, network structure, interactions between the wallets and underlying business processes. Also design problems can be found and potential privacy and security issues.

The network structure reveals information about the major accounts and their activity. Due to these findings we were able to discover, the top 5 receivers and senders in the network.

As the data that shows transactions varying from one day to another, from a very low level to relatively high transaction volume. This gave the impression that there are major nodes which distribute disproportional amounts of Ether in the network. The data showed, that the Top 5 senders perform 36 percent of the transactions in the network. This discrepancy could be shown between senders and receivers during the examined period. We surmise, that the 4 main receivers belong to currency exchanges. All of them appeared to be smart contracts, same as the first major receiver. Through online research, we were able to identify that wallet as a smart contract designed to deal with the side effects of the hard fork in 2016. Due to the fact, that there are major nodes one could identify indirect the underlying business processes. The discrepancies in the transaction volume could not be investigated conclusively just through the scraped data alone. But a side research showed that the top 5 transaction senders can be identified as mining pools.

An alternative hypothesis, which would involve potential privacy and security issues could neither be confirmed nor disproved.

To find some major nodes in the Ethereum network, the distribution of several centrality measures were computed and interpreted. Major results are, that within the network one can find major hubs, with high transaction volume. Most of the nodes are not connected. They are usually just having little transactions between each other. The fit with a theoretical power law distribution shows, the *Small World Phenomenon* can only be confirmed for subgraphs within the network. This is also expected given the fact that the graph analysis found out that there a very little nodes which cover moderate to high amounts of transactions.

Originally the Ethereum protocol was thought of as a modified version of the Bitcoin Blockchain. That's what Vitalik Buterin had in mind when he wrote the white paper in 2014. The Turing completeness should provide opportunities to create applications for any type of application [Buterik, 2013].

Vitalik's idea was to optimize the Ethereum Blockchain technology to make the blockchain technology more applicable for a larger range of applications. He thought Ethereum being open ended by design and believes that it is extremely well suited to serve as a basic layer for financial and non-financial

applications [Buterik, 2013].

As shown in the comparison between blockchain and bitcoin, one could guess, that so far this goal is not close to be reached by the Ethereum community. There are lots of attempts to create applications for the platform and try to find business cases. An indication for that claim is the fact that the number of actively used Ethereum addresses has dropped below 1 million and also that a study on usage of smart contracts indicated, that 95% are used just less than 10 times and only 5% more than that [Chandersekhar, 2018]. An interesting question would be, how many contracts exist, that are just using up computational capacity without being used more than a few times.

To conclude, we were not able to perform such a comprehensive study as it was done from Lischke and Fabian and the insights into the Ethereum transactions and regarding information are less than expected. We had some limitations along the way.

A crucially limiting factor was the access to the different data sets and the amount of data we were able to collect. The data was incomplete, and we were not able to locate the transactions via IP addresses. In the replicated paper they were able to scrape the data with the transaction data, using 'blockchain.info'. This is not possible for the Ethereum blockchain and a proper workaround was not found easily. We searched for alternative providers, but none of them gave information about the sender's IP address. Thus, we could not trace the sender's geographical location, or find a mapping to the wallet's industry tag either. Another challenge is, that the network is, in essence, different from the Bitcoin. Especially, because the Ethereum blockchain also provides smart contracts and a currency service. Where the distinction is not easily made in the data. Some use the platform just for transactions and a lot of them apply smart contracts or decentralized applications.

As mentioned in the methods, for a further paper one could look in more detail into the scraper we applied and modify it for the usage in the 'go' language.

To get back to the data, unfortunately the time period gathered was during a time when a hard fork happened in the Ethereum network, in 2017. All that happened because of a hard fork incidence, when a complicated hack of a smart contracts led to a loss of 12,7 million US Dollar. The network got divided and had to reorganize. This brought noise in the data, which is not the usual situation and to get broader explanations, one should take care to use a different period or a longer one, to get more general insights. However, one is also hard pressed to find a period in the Ethereum blockchain that does not experience heavy turmoil.

## VIII. APPENDIX



### 1) Configuration:

```
1000 from configparser import ConfigParser
1002
1004 def config(filename="database.ini", section="postgresql"):
1006     parser = ConfigParser()
1008     parser.read(filename)
1010
1012     db = {}
1014     if parser.has_section(section):
1016         params = parser.items(section)
1018         for param in params:
1020             db[param[0]] = param[1]
1022     else:
1024         msg = "Section {} not found in the {} file"
1026         raise Exception(msg.format(section, filename))
1028
1030     return db
```

../preprocessing/database/config.py

```
1000 import os
1002 import psycopg2
1004 from config import config
1006 from glob import glob
1008
1010 def create_tables(command):
1012     """Create tables for ethereum database."""
1014     conn = None
1016     try:
1018         params = config()
1020         conn = psycopg2.connect(**params)
1022         cur = conn.cursor()
1024         cur.execute(command)
1026         cur.close()
1028         # If you don't commit, no changes will be made.
1030         conn.commit()
1032     except (Exception, psycopg2.DatabaseError) as error:
1034         print(error)
1036     finally:
1038         if conn is not None:
1040             conn.close()
1042
1044 if __name__ == "__main__":
1046     queries_path = os.path.join("sql", "create", "*")
1048     queries = glob(queries_path)
1050     # HACK: table ip-info depends on transactoins, sort to fix order.
1052     queries.sort(reverse=True)
1054     for query in queries:
1056         with open(query) as q:
1058             command = q.read()
1060             create_tables(command)
```

../preprocessing/database/create\_tables.py

### A. Create Tables

```
1000 CREATE TABLE IF NOT EXISTS blocks (
1002     block_hash VARCHAR(255) PRIMARY KEY,
1004     gas_used INTEGER,
1006     gas_limit INTEGER,
1008     block_id INTEGER UNIQUE,
1010     inception_time INTEGER
1012 )
```

../preprocessing/database/sql/create/block.sql

```
1000 CREATE TABLE IF NOT EXISTS ip_info (
1002     transaction_hash VARCHAR(255) NOT NULL,
1004     ip_address VARCHAR(255) NOT NULL
```

```
)
```

../preprocessing/database/sql/create/ip\_info.sql

```
1000 CREATE TABLE IF NOT EXISTS trade (
1001     trade_date INTEGER PRIMARY KEY,
1002     open FLOAT,
1003     high FLOAT,
1004     low FLOAT,
1005     close FLOAT,
1006     value_eth FLOAT,
1007     value_usd FLOAT
1008 )
```

../preprocessing/database/sql/create/trade.sql

```
1000 CREATE TABLE IF NOT EXISTS transactions (
1001     transaction_hash VARCHAR(255) NOT NULL,
1002     block_hash VARCHAR(255) NOT NULL,
1003     sender VARCHAR(255),
1004     receiver VARCHAR(255),
1005     value FLOAT,
1006     gas_used INTEGER,
1007     gas_price BIGINT
1008 )
```

../preprocessing/database/sql/create/transaction.sql

```
1000 CREATE TABLE IF NOT EXISTS users (
1001     user_id SERIAL PRIMARY KEY,
1002     wallet_address VARCHAR(255) NOT NULL
1003 )
```

../preprocessing/database/sql/create/users.sql

## B. ETL: Extract Transfer Load

```
1000 import psycopg2
1001 from config import config
1002
1003 def connect():
1004     """Connect to the PostgreSQL database server."""
1005     conn = None
1006     try:
1007         params = config()
1008         # Create connection.
1009         print("Connecting to the PostgreSQL database...")
1010         conn = psycopg2.connect(**params)
1011         # Create cursor.
1012         cur = conn.cursor()
1013
1014         # Execute a statement
1015         cur.execute("SELECT version()")
1016
1017         # Fetch results and display them.
1018         db_version = cur.fetchone()
1019         print("Database version:", db_version)
1020
1021         # Close connections
1022         cur.close()
1023
1024     except (Exception, psycopg2.DatabaseError) as error:
1025         print(error)
1026
1027     finally:
1028         if conn is not None:
1029             conn.close()
1030             print("Database connection closed.")
1031
1032
```

```
1034 if __name__ == "__main__":
    connect()
```

../preprocessing/database/connect.py

```
1000 from configparser import ConfigParser
1002
1003 def config(filename="database.ini", section="postgresql"):
1004     parser = ConfigParser()
1005     parser.read(filename)
1006
1007     db = {}
1008     if parser.has_section(section):
1009         params = parser.items(section)
1010         for param in params:
1011             db[param[0]] = param[1]
1012     else:
1013         msg = "Section {} not found in the {} file"
1014         raise Exception(msg.format(section, filename))
1016
1017     return db
```

../preprocessing/database/config.py

```
1000 import os
1001 import psycopg2
1002 from config import config
1003 from web3.auto import w3
1004 from hexbytes import HexBytes
1005 from blockchain import blockexplorer
1006
1007 def extract_info(dictlike, keys):
1008     """
1009     Extract values from a dictionary-style object.
1010
1011     * Parameters:
1012         * 'dictlike': object that has all attributes in 'keys'
1013         * 'keys': list of keys to loop over.
1014
1015     * Returns:
1016         * 'out': tuple of values corresponding to each key in 'keys'.
1017     """
1018     out = []
1019
1020     for key in keys:
1021         val = getattr(dictlike, key, "NULL")
1022         if val is None:
1023             val = "NULL"
1024         if type(val) == HexBytes:
1025             out.append(val.hex())
1026         else:
1027             out.append(val)
1028
1029     return tuple(out)
1030
1031 def update(dictlike, table, keys):
1032     """
1033     Update a database-table with the contents of one block.
1034
1035     * Parameters:
1036         * 'dictlike': object that has a '__get__'-method implemented.
1037         * 'table': ('str') name of table to update.
1038         * 'keys': list of keys to loop over.
1039     """
1040
1041     row = extract_info(dictlike, keys)
1042
1043     sql = "INSERT INTO {} VALUES {}".format(table, row)
1044     conn = None
1045     try:
```

```

1048     params = config()
1049     conn = psycopg2.connect(**params)
1050     cur = conn.cursor()
1051     cur.execute(sql)
1052     conn.commit()
1053     cur.close()
1054 except (Exception, psycopg2.DatabaseError) as error:
1055     print(error)
1056 finally:
1057     if conn is not None:
1058         conn.close()
1059
1060 def update_list(block, table, keys, block_key="transactions"):
1061     """
1062     Batch update a database-table with the contents of one block.
1063
1064     * Parameters:
1065     * 'block': object that has a '__get__'-method implemented.
1066     * 'table': ('str') name of table to update.
1067     * 'keys': list of keys to loop over.
1068     * 'block_key': ('str') key of 'block' that contains an iterable.
1069     """
1070     conn = None
1071     try:
1072         params = config()
1073         conn = psycopg2.connect(**params)
1074         cur = conn.cursor()
1075
1076         for transaction in block[block_key]:
1077             tx = w3.eth.getTransaction(transaction)
1078
1079             info = extract_info(tx, keys)
1080             sql = "INSERT INTO {} VALUES {}".format(table, info)
1081             cur.execute(sql)
1082             conn.commit()
1083             cur.close()
1084     except (Exception, psycopg2.DatabaseError) as error:
1085         print(error)
1086     finally:
1087         if conn is not None:
1088             conn.close()
1089
1090
1091 def tx_iterator(block):
1092     for transaction in block.transactions:
1093         yield transaction.hex()
1094
1095
1096 if __name__ == "__main__":
1097     block = w3.eth.getBlock(6753133)
1098     block_keys = ["hash", "gasUsed", "gasLimit", "number", "timestamp"]
1099     tx_keys = ["blockHash", "hash", "from", "to", "value", "gas", "gasPrice"]
1100     ip_keys = ["hash", "relayed_by"]
1101
1102     update(block, 'blocks', block_keys)
1103     update_list(block, 'transactions', tx_keys)
1104     for tx in tx_iterator(block):
1105         try:
1106             transaction = blockexplorer.get_tx(tx)
1107             update(transaction, 'ip_info', ip_keys)
1108         except Exception as e:
1109             print(e)
1110             continue
1111
1112
1113 def query(block):
1114     with open(os.path.join("sql", "query", "extract_info.sql")) as q:
1115         sql = q.read()
1116     conn = None
1117     try:
1118         params = config()
1119         conn = psycopg2.connect(**params)
1120         cur = conn.cursor()

```

```

1122         cur.execute(sql.format(block))
1123         rows = cur.fetchall()
1124         cur.close()
1125     except (Exception, psycopg2.DatabaseError) as error:
1126         print(error)
1127     finally:
1128         if conn is not None:
1129             conn.close()
1130
1131     return rows

```

../preprocessing/database/database.py

```

1000 import os
1001 import sys
1002 import database
1003 import logging
1004 from web3.auto import w3
1005 from web3.exceptions import UnhandledRequest
1006 from tqdm import trange
1007 from blockchain import blockexplorer, exceptions
1008
1009 # Read api-key
1010 if os.path.exists("api_code.txt"):
1011     with open("api_code.txt") as f:
1012         api_code = f.read().strip()
1013 else:
1014     api_code = None
1015
1016 # Get logger.
1017 logger = logging.getLogger(__name__)
1018 logger.setLevel(logging.INFO)
1019 fh = logging.FileHandler("build.log")
1020 format = '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
1021 formatter = logging.Formatter(format)
1022 fh.setFormatter(formatter)
1023 logger.addHandler(fh)
1024
1025 if __name__ == "__main__":
1026     # General Setup at first start.
1027     if not os.path.exists("progress.txt"):
1028         logger.info("Start ETL.")
1029         broken_file = open("broken.txt", "w")
1030         broken_file.close()
1031         progress_file = open("progress.txt", "w")
1032         progress_file.close()
1033         # First block in chain.
1034         start_block_nr = 2000000
1035     # Resume from broken pipe.
1036     else:
1037         if os.stat("progress.txt").st_size == 0:
1038             start_block_nr = 2000000
1039         else:
1040             with open("progress.txt", "r") as f:
1041                 # The block to continue from, is the last line in the file.
1042                 for line in f:
1043                     pass
1044
1045             start_block_nr = int(line)
1046             logger.info("Resume ETL at {}".format(start_block_nr))
1047
1048     # Get latest block and associated number.
1049     end_block = w3.eth.getBlock("latest")
1050     end_block_nr = end_block["number"]
1051
1052     # Fix: sometimes the latest block is unconfirmed and it's number is 0
1053     if end_block_nr == 0:
1054         end_block_nr = start_block_nr + 1000000
1055
1056     block_keys = ["hash", "gasUsed", "gasLimit", "number", "timestamp"]
1057     tx_keys = ["hash", "blockHash", "from", "to", "value", "gas", "gasPrice"]
1058     ip_keys = ["hash", "relayed_by"]
1059
1060     for current_block_nr in trange(start_block_nr, end_block_nr+1):

```

```

1062     try:
1063         block = w3.eth.getBlock(current_block_nr)
1064         database.update(block, 'blocks', block_keys)
1065         database.update_list(block, 'transactions', tx_keys)
1066         for tx in database.tx_iterator(block):
1067             try:
1068                 transaction = blockexplorer.get_tx(tx, api_code=api_code)
1069                 database.update(transaction, 'ip_info', ip_keys)
1070             except exceptions.APIException:
1071                 continue
1072             except Exception as e:
1073                 logger.error(e)
1074                 continue
1075             # Save progress
1076             with open("progress.txt", "a") as f:
1077                 f.write(str(current_block_nr) + "\n")
1078         except UnhandledRequest as error:
1079             logger.error(error)
1080             sys.exit(1)
1081
1082         except (EOFError, KeyboardInterrupt):
1083             logger.info("Aborting ETL")
1084             sys.exit(0)
1085         except Exception as error:
1086             logger.error(error)
1087             with open("broken.txt", "a") as f:
1088                 f.write(str(current_block_nr) + "\n")
1089             continue
1090
1091     logger.info("End of ETL.")

```

../preprocessing/database/build.py

### C. Query Database

```

1000 import os
1001 import logging
1002 from tqdm import trange
1003 import database
1004 import pandas as pd
1005
1006 # Get logger.
1007 logger = logging.getLogger(__name__)
1008 logger.setLevel(logging.INFO)
1009 fh = logging.FileHandler("query.log")
1010 format = '%(asctime)s - %(name)s - %(levelname)s - %(message)s'
1011 formatter = logging.Formatter(format)
1012 fh.setFormatter(formatter)
1013 logger.addHandler(fh)
1014
1015 datafile = os.path.join(".", "..", "data", "ethereum-data.csv")
1016 columns = ["block_hash", "block_gas", "gas_limit", "inception_time", "tx_hash",
1017            "sender", "receiver", "value", "gas_used", "gas_price"]
1018
1019 for block in trange(2000000, 2014806):
1020     try:
1021         blockinfo = database.query(block)
1022         df = pd.DataFrame(blockinfo, columns=columns)
1023         if os.path.exists(datafile):
1024             with open(datafile, "a") as f:
1025                 df.to_csv(f, header=False, index=False)
1026         else:
1027             df.to_csv(datafile, index=False)
1028     except Exception as e:
1029         logger.error(e)
1030

```

../preprocessing/database/query.py

```

1000 import os
1001 import psycpg2
1002 import pandas as pd

```



```

1004 from config import config

1006
1008 datafile = os.path.join("../", "../", "data", "ethereum-data-large.csv")
1008 columns = ["block_hash", "block_gas", "gas_limit", "inception_time", "tx_hash",
1010            "sender", "receiver", "value", "gas_used", "gas_price"]
1010
1012 with open(os.path.join("sql", "query", "extract_batch.sql")) as q:
1012     sql = q.read()
1012     conn = None
1014     try:
1016         params = config()
1016         conn = psycopg2.connect(**params)
1016         cur = conn.cursor()
1018         cur.execute(sql)
1018         rows = cur.fetchall()
1020         cur.close()
1022     except (Exception, psycopg2.DatabaseError) as error:
1022         print(error)
1022     finally:
1024         if conn is not None:
1024             conn.close()
1026
1028 df = pd.DataFrame(rows, columns=columns)
1028 df.to_csv(datafile, index=False)

```

../preprocessing/database/batch\_query.py

```

1000 SELECT
1002     b.block_hash ,
1002     b.block_gas ,
1004     b.gas_limit ,
1004     b.inception_time ,
1006     t.transaction_hash ,
1006     t.sender ,
1008     t.receiver ,
1008     t.value ,
1010     t.gas_used ,
1010     t.gas_price
1012 FROM (
1012     SELECT
1014         block_hash ,
1014         gas_used AS block_gas ,
1016         gas_limit ,
1016         inception_time
1018     FROM blocks) b
1018 LEFT JOIN (
1020     SELECT
1020         transaction_hash ,
1022         block_hash ,
1022         sender ,
1024         receiver ,
1024         value ,
1026         gas_used ,
1026         gas_price
1028     FROM transactions) t
1028 ON t.block_hash = b.block_hash

```

../preprocessing/database/sql/query/extract\_batch.sql

```

1000 WITH blocktable AS (
1002     SELECT
1002         block_hash ,
1004         gas_used AS block_gas ,
1004         gas_limit ,
1006         inception_time
1006     FROM blocks
1008     WHERE block_id = {}
1008 ),
1010 transactiontable AS (

```

```

1012     SELECT
1013         transaction_hash ,
1014         block_hash ,
1015         sender ,
1016         receiver ,
1017         value ,
1018         gas_used ,
1019         gas_price
1020     FROM transactions
1021 )
1022
1023 SELECT *
1024 FROM blocktable bt
1025 LEFT JOIN transactiontable tt USING(block_hash);

```

../preprocessing/database/sql/query/extract\_info.sql

#### D. Analysis

```

1000 import os
1001 import powerlaw
1002 import networkx as nx
1003 import numpy as np
1004 import pandas as pd
1005 from matplotlib import pyplot as plt
1006 from matplotlib.ticker import FormatStrFormatter
1007
1008 data_path = os.path.join(".", "data", "ethereum-data.zip")
1009 general_data_path = os.path.join(".", "data", "export-EtherPrice.csv")
1010 outdir = os.path.join(".", "analysis")
1011
1012 # Load custom dataset
1013 data = pd.read_csv(data_path)
1014 data.inception_time = pd.to_datetime(data.inception_time ,
1015                                     unit="s")
1016
1017 data["day"] = data.inception_time.dt.date
1018 # Convert to ether
1019 data["value"] = data["value"] / 10**18
1020
1021 # Load general dataset
1022 general_data = pd.read_csv(general_data_path ,
1023                           names=["date", "timestamp", "val"], skiprows=[0])
1024 general_data.date = pd.to_datetime(general_data.date)
1025 general_data["day"] = general_data.date.dt.date
1026
1027 # Merge dataframes
1028 data = pd.merge(data, general_data, how="left", on="day")
1029 data.columns
1030 data.agg({"block_hash": "nunique", "tx_hash": "nunique"})
1031
1032 senders = set(data.sender)
1033 receivers = set(data.receiver)
1034
1035 tx_dist = data.sender.value_counts().value_counts(normalize=True)[:10]
1036 breaks = tx_dist.iloc[[0, 2, 3, 9]]
1037
1038 fig, ax = plt.subplots(1)
1039 tx_dist.plot.bar(color="red", ax=ax, rot=0)
1040 ax.set_ylabel("Percentage of wallets")
1041 ax.set_xlabel("Total transactions")
1042 ax.set_yticks(breaks)
1043 ax.yaxis.set_major_formatter(FormatStrFormatter('%.2f'))
1044 ax.hlines(breaks, xmin=-10, xmax=10, color="white", linewidth=1, alpha=1)
1045 ax.set_xlim([-0.5, 9.5])
1046 plt.tight_layout()
1047 fig.savefig(os.path.join("pics", "distribution.png"))
1048 plt.clf()
1049 # Descriptive statistics for custom dataset
1050 desc_cols = ["sender", "receiver", "value", "gas_used", "gas_price"]
1051 desc = data[desc_cols].describe().round(2).iloc[1:, :]
1052 desc.columns = ["value", "gas used", "gas price"]

```

```

1054 # Round data and add num users/receivers/senders
desc.to_latex(os.path.join(outdir, "descriptive.tex"))
1056 desc.to_csv(os.path.join(outdir, "descriptive.csv"), index=False)

1058 daily = data.groupby("day").aggregate({"sender": "nunique",
1060                                         "receiver": "nunique",
1062                                         "value": [np.mean, np.median, np.std],
1064                                         "val": "mean"})

receivers = data.loc[:, ["receiver", "day"]]
1066 receivers.columns = ["user", "day"]
senders = data.loc[:, ["sender", "day"]]
senders.columns = ["user", "day"]

1068 users = pd.concat([receivers, senders])
daily["users"] = users.groupby("day").user.nunique()
1070 daily = daily[["sender", "receiver", "users", "value", "val"]]

1072 # =====
# Draw the graph
1074 sender = set(data.sender.values)
receiver = set(data.receiver.values)
1076 interactors = sender.intersection(receiver)

1078 loc = data.sender.isin(interactors) & data.receiver.isin(interactors)
interactor_data = data.loc[loc, ["sender", "receiver", "value"]]
1080 interactor_data.shape
len(interactors)

1082 fig, ax = plt.subplots(1, figsize=(15, 15))
1084 sample = interactor_data.sample(2000)
G = nx.from_pandas_edgelist(df=sample,
1086                           source="sender",
1088                           target="receiver",
1090                           edge_attr="value")
G.add_nodes_from(nodes_for_adding=sample.sender.tolist())
1092 nx.draw(G, with_labels=False, ax=ax)
fig.savefig(os.path.join(outdir, "graph-subsample-1000.png"))

1094 fig, ax = plt.subplots(1, figsize=(20, 20))
G = nx.from_pandas_edgelist(df=interactor_data,
1096                           source="sender",
1098                           target="receiver",
1100                           edge_attr="value")
G.add_nodes_from(nodes_for_adding=interactor_data.sender.tolist())
1102 nx.draw(G, with_labels=False, ax=ax)
fig.savefig(os.path.join(outdir, "graph-full.png"))

1104 # Calculate Network metrics
idx = pd.Index(interactors)
1106 outgoing = interactor_data.sender.value_counts().reindex(idx, fill_value=0)
ingoing = interactor_data.receiver.value_counts().reindex(idx, fill_value=0)

1108 edges = pd.concat([outgoing.rename("outgoing"), ingoing.rename("ingoing")],
1110                   axis=1, levels=["ingoing", "outgoing"])
edges["total"] = outgoing + ingoing

# Centrality
1112 triangles = nx.triangles(G)
edges["triangles"] = pd.Series(triangles).reindex(idx, fill_value=0)
1114 edges["C_in"] = 2*edges["triangles"] / (edges["ingoing"]*(edges["ingoing"]-1))
edges["C_out"] = 2*edges["triangles"] / (edges["outgoing"]*(edges["outgoing"]-1))
1116 edges["C_total"] = 2*edges["triangles"] / (edges["total"]*(edges["total"]-1))

edges.replace([np.inf, np.nan], 0, inplace=True)
1118 avg_cluster = edges[["C_in", "C_out", "C_total"]].mean()

1120 edges.C_in.plot(kind="hist")
1122 edges.C_out.plot(kind="density", ylim=(0, 1))
edges.C_total.plot(kind="density")

1124 # Fit power law
1126 fig, ax = plt.subplots(1)

```

```

1128 def plot_dist(data, ax):
1130     """Compare theoretical and empirical powerlaw distribution."""
1132     data = data[data > 0]
1134     fit = powerlaw.Fit(data, discrete=True)
1136     fit.plot_pdf(color='b', linewidth=2, ax=ax, label="PDF")
1138     fit.power_law.plot_pdf(color='b', linestyle="—", ax=ax,
1140                           label="Power Law fit PDF")
1142     fit.plot_ccdf(color='r', linewidth=2, ax=ax, label="CCDF")
1144     fit.power_law.plot_ccdf(color='r', linestyle="—", ax=ax,
1146                             label="Power Law fit CCDF")
1148     ax.set_ylabel(r"p")
1150     ax.legend(loc="lower left")
1152
1154 fig, axes = plt.subplots(ncols=3, figsize=(20, 5))
1156 titles = ["Transactions sent", "Transactions received", "Total transactions"]
1158 datalist = [edges.outgoing, edges.ingoing, edges.total]
1160 for title, ax, data in zip(titles, axes, datalist):
1162     plot_dist(data.values, ax)
1164     ax.set_title(title)
1166
1168 fig.savefig(os.path.join(outdir, "power-law-fit.png"))
1170
1172 # Playground Area
1174 # Degree centrality
1176 deg_cent = nx.degree_centrality(G)
1178 edges["deg_centrality"] = pd.Series(deg_cent).reindex(idx)
1180
1182 # Between centrality
1184 bet_cent = nx.betweenness_centrality(G)
1186 edges["betw_centrality"] = pd.Series(bet_cent).reindex(idx)
1188 avg_cluster
1190 avg_cluster.to_latex()
1192
1194 deg_cent = nx.degree_centrality(G)
1196 bet_deg_cent = nx.betweenness_centrality(G)
1198 closeness_cent = nx.closeness_centrality(G)
1200 node_clustering = nx.clustering(G)
1202
1204 avg_clustering = nx.average_clustering(G)
1206
1208 centrality = pd.DataFrame({
1210     "degree centrality": deg_cent,
1212     "betweenness centrality": bet_deg_cent,
1214     "closeness centrality": closeness_cent,
1216     "local clustering": node_clustering
1218 })

```

../analysis/analysis.py

```

1000 """
1002 Create figures for Centrality Measures.
1004 """
1006 import os
1008 import pandas as pd
1010 from matplotlib import pyplot as plt
1012
1014 data_path = os.path.join("analysis", "centrality_100000.csv")
1016
1018 data = pd.read_csv(data_path, index_col=0)
1020
1022 data.shape
1024 data.columns = data.columns.str.capitalize()
1026 desc = data.describe().iloc[1:, :].round(4)
1028 desc.to_latex(os.path.join("analysis", "centrality.tex"))
1030
1032 fig, axes = plt.subplots(2, 2)
1034 for col, ax in zip(data.columns, axes.flatten()):
1036     data[col].plot.box(ax=ax)

```

```
1020 fig.savefig(os.path.join("analysis", "centrality.png"))
```

../analysis/centrality\_figures.py

```
1000 import os
1001 import networkx as nx
1002 import pandas as pd
1003 from sys import argv
1004
1005 _, sample_size = argv
1006 data_path = os.path.join(".", "data", "ethereum-data.zip")
1007 general_data_path = os.path.join(".", "data", "export-EtherPrice.csv")
1008 outdir = os.path.join(".", "analysis")
1009
1010 data = pd.read_csv(data_path)
1011 data.inception_time = pd.to_datetime(data.inception_time,
1012                                     unit="s")
1013
1014 data["day"] = data.inception_time.dt.date
1015 # Convert to ether
1016 data["value"] = data["value"] / 10**18
1017
1018 # Load general dataset
1019 general_data = pd.read_csv(general_data_path,
1020                            names=["date", "timestamp", "val"], skiprows=[0])
1021 general_data.date = pd.to_datetime(general_data.date)
1022 general_data["day"] = general_data.date.dt.date
1023
1024 # Merge dataframes
1025 data = pd.merge(data, general_data, how="left", on="day")
1026 receivers = data.loc[:, ["receiver", "day"]]
1027 receivers.columns = ["user", "day"]
1028 senders = data.loc[:, ["sender", "day"]]
1029 senders.columns = ["user", "day"]
1030
1031 # =====
1032 # Draw the graph
1033 sender = set(data.sender.values)
1034 receiver = set(data.receiver.values)
1035 interactors = sender.intersection(receiver)
1036
1037 loc = data.sender.isin(interactors) & data.receiver.isin(interactors)
1038 interactor_data = data.loc[loc, ["sender", "receiver", "value"]]
1039
1040 print("Initialize graph")
1041 sample = interactor_data.sample(10000)
1042 G = nx.from_pandas_edgelist(df=sample,
1043                             source="sender",
1044                             target="receiver",
1045                             edge_attr="value")
1046 G.add_nodes_from(nodes_for_adding=sample.sender.tolist())
1047
1048 print("Compute Centrality-Measures")
1049 deg_cent = nx.degree_centrality(G)
1050 bet_deg_cent = nx.betweenness_centrality(G)
1051 closeness_cent = nx.closeness_centrality(G)
1052 node_clustering = nx.clustering(G)
1053
1054 avg_clustering = nx.average_clustering(G)
1055 print("Average clustering coefficient:", avg_clustering)
1056
1057 centrality = pd.DataFrame({
1058     "degree centrality": deg_cent,
1059     "betweenness centrality": bet_deg_cent,
1060     "closeness centrality": closeness_cent,
1061     "local clustering": node_clustering
1062 })
1063
1064 print("Save to .csv")
1065 centrality.to_csv(os.path.join(outdir, "centrality_"+str(sample_size)+".csv"))
```

../analysis/centrality.py

```

1000 if(!require("data.table")) install.packages("data.table")
1001 if(!require("dplyr")) install.packages("dplyr")
1002 if(!require("ggplot2")) install.packages("ggplot2")
1003 if(!require("lubridate")) install.packages("lubridate")
1004 if(!require("maptools")) install.packages("maptools")
1005 if(!require("scales")) install.packages("scales")
1006
1007 library("data.table")
1008 library("dplyr")
1009 library("ggplot2")
1010 library("lubridate")
1011 library("maptools")
1012 library("scales")
1013 library(xtable)
1014
1015 # Data transformations
1016 e_w_rate <- 1000000000000000000 #1 Ether equals x wei
1017 s_w_rate <- 1000000000000 #1 Szabo equals x wei
1018
1019 dt <- read.csv("ethereum-data-large.csv", colClasses = "character")
1020 s1 <- summary(dt)
1021
1022 colnames_1 <- c("block_hash", "tx_hash", "sender", "receiver")
1023 colnames_2 <- c("block_gas", "gas_limit", "inception_time",
1024               "value", "gas_used", "gas_price")
1025 dt[colnames_1] <- lapply(dt[colnames_1], as.factor)
1026 dt[colnames_2] <- lapply(dt[colnames_2], as.numeric)
1027
1028 dt$inception_time <- as.POSIXct(dt$inception_time, origin="1970-01-01")
1029 dt$time <- trunc(dt$inception_time, units = "hours")
1030 dt$time <- as.POSIXct(dt$time, origin="1970-01-01")
1031 dt$day <- trunc(dt$inception_time, units = "days")
1032 dt$day <- as.POSIXct(dt$day, origin="1970-01-01")
1033 s2 <- summary(dt)
1034
1035 dt <- dt[!is.na(dt$value),]
1036
1037 dt$fee <- dt$gas_used*dt$gas_price
1038
1039 dt_sum <- lapply(dt[,c("value", "gas_used", "gas_price", "fee")], sum)
1040 dt_median <- lapply(dt[,c("value", "gas_used", "gas_price", "fee")], median)
1041 dt_max <- lapply(dt[,c("value", "gas_used", "gas_price", "fee")], max)
1042
1043 #——Data description——
1044 s3 <- summary(dt)
1045
1046 # Graphs
1047
1048 hist(dt$gas_used, col = "orange", border = "grey",
1049      main = "Distribution of Gas used per Transaction",
1050      xlab = "Amount of Gas per Transaction, in wei",
1051      ylab = "Number of Transactions", ylim = c(0,70000))
1052 boxplot(dt$gas_used,
1053        # main = "Boxplot of Gas used per Transaction",
1054        xlab = "Amount of Gas per Transaction, in wei")
1055
1056 # hist(dt$gas_price)
1057 # boxplot(dt$gas_price)
1058
1059 hist(dt$fee/e_w_rate, col = "orange", border = "grey",
1060      # main = "Distribution of Fees paid per Transaction",
1061      xlab = "Fee per Transaction, in ether",
1062      ylab = "Number of Transactions",
1063      ylim = c(0,700),
1064      xlim = c(0,2),
1065      breaks = 25000)
1066 boxplot(dt$fee/e_w_rate,
1067        xlab = "Fee per Transaction, in ether")
1068
1069 # CAREFULL!! THIS TAKES TX to make a calculation
1070 hist(dt$value/e_w_rate, col = "orange", border = "grey",
1071      xlab = "Value of a Transaction, in ether",
1072      ylab = "Number of Transactions",

```



```

1074     ylim = c(0,50),
        breaks = 50)
1076 boxplot(dt$value/e_w_rate,
          xlab = "Value of a Transaction, in ether")

1078 # Add fee calculation
tx_n <- nrow(dt)
1080 tx_volume <- sum(dt$value, na.rm = TRUE)
tx_average_value <- tx_n/tx_volume
1082 tx_n; tx_volume; tx_average_value

1084 block_descriptives <- dt %>%
  group_by(block_hash) %>%
1086   summarize(block_tx=n(),
              block_tx_perc=n()/tx_n*100,
1088             block_volume=sum(value),
              block_volume_perc=sum(value)/tx_volume*100,
1090             block_average_tx_value=mean(value)) %>%
  arrange(desc(block_tx))
1092 summary(block_descriptives)

1094 dt_sum$block_tx <- sum(block_descriptives$block_tx)
dt_sum$block_volume <- sum(block_descriptives$block_volume)
1096
1098 dt_median$block_tx <- median(block_descriptives$block_tx)
dt_median$block_volume <- median(block_descriptives$block_volume)

1100 dt_max$block_tx <- max(block_descriptives$block_tx)
dt_max$block_volume <- max(block_descriptives$block_volume)
1102
1104 hist(block_descriptives$block_tx, col = "orange", border = "grey",
      xlab = "Number of Transactions",
      ylab = "Number of Blocks",
1106      ylim = c(0,7000))
1108 hist(block_descriptives$block_volume/e_w_rate,
      col = "orange", border = "grey",
      xlab = "Volume of Transactions, in ether",
1110      ylab = "Number of Blocks", ylim = c(0,100), breaks = 100)

1112 time_descriptives <- dt %>%
  group_by(day) %>%
1114   summarize(tx_number=n(),
              tx_perc=n()/tx_n*100,
1116             tx_vol=sum(value, na.rm = TRUE),
              tx_vol_perc=sum(value, na.rm = TRUE)/tx_volume*100,
1118             tx_average_value=mean(value, na.rm = TRUE)) %>%
  arrange(desc(day))
1120
1122 # remove first and last day which are not complete
time_descriptives <- time_descriptives[-c(1, nrow(time_descriptives)), ]
summary(time_descriptives)
1124
1126 dt_sum$tx_number <- sum(time_descriptives$tx_number)
dt_sum$tx_volume <- sum(time_descriptives$tx_vol)
1128
1130 dt_median$tx_number <- median(time_descriptives$tx_number)
dt_median$tx_volume <- median(time_descriptives$tx_vol)

1132 dt_max$tx_number <- max(time_descriptives$tx_number)
dt_max$tx_volume <- max(time_descriptives$tx_vol)
1134
1136 ggplot(data=time_descriptives, aes(x=day, y=tx_number)) +
  geom_line(colour="blue") +
  ylab("Number of Transactions") +
1138  xlab("Date")

1140 ggplot(data=time_descriptives, aes(x=day, y=tx_vol/e_w_rate)) +
  geom_line(colour="blue") +
1142  # ggtitle("Volume of Transactions over Time, in ether") +
  ylab("Volume of Transactions") +
1144  xlab("Date")
1146

```

```

# Major wallets
1148 major_senders <- dt %>%
  group_by(sender) %>%
1150   summarize(sender_tx=n(),
               sender_tx_perc=n()/tx_n*100,
1152   sender_volume=sum(value),
               sender_volume_perc=sum(value)/tx_volume*100,
1154   sender_average_value=mean(value)) %>%
  arrange(desc(sender_tx))
1156 colnames(major_senders)[1] <- "wallet"
summary(major_senders)

1158
major_receivers <- dt %>%
1160   group_by(receiver) %>%
  summarize(receiver_tx=n(),
1162   receiver_tx_perc=n()/tx_n*100,
               receiver_volume=sum(value),
1164   receiver_volume_perc=sum(value)/tx_volume*100,
               receiver_average_value=mean(value)) %>%
1166   arrange(desc(receiver_tx))
colnames(major_receivers)[1] <- "wallet"
1168 summary(major_receivers)

1170 major_senders <- data.table(major_senders)
major_receivers <- data.table(major_receivers)
1172 setkey(major_receivers, wallet)
setkey(major_senders, wallet)
1174 m_w <- merge(major_senders, major_receivers, all = TRUE)
m_w[is.na(m_w)] <- 0
1176 m_w$number_tx <- m_w$sender_tx + m_w$receiver_tx
m_w$vol_tx <- m_w$sender_volume + m_w$receiver_volume
1178 m_w$balance <- m_w$sender_volume + m_w$receiver_volume
summary(m_w)

1180
1182 top50_sender_tx <- head(m_w[order(-sender_tx),c("wallet", "sender_tx")], 50)
top50_sender_tx$cum <- cumsum(top50_sender_tx$sender_tx)
1184 top50_sender_tx$perc <- top50_sender_tx$cum/tx_n*100
plot(top50_sender_tx$perc, ylim = c(0,100),
1186   xlab = "Top 50 Senders",
   ylab = "Percentage of all Transactions Number")
1188
top50_receiver_tx <- head(m_w[order(-receiver_tx),c("wallet", "receiver_tx")], 50)
top50_receiver_tx$cum <- cumsum(top50_receiver_tx$receiver_tx)
1190 top50_receiver_tx$perc <- top50_receiver_tx$cum/tx_n*100
plot(top50_receiver_tx$perc, ylim = c(0,100),
1192   xlab = "Top 50 Receivers",
   ylab = "Percentage of all Transactions Number")
1194
1196
top50_sender_volume <- head(m_w[order(sender_volume),c("wallet", "sender_volume")], 50)
top50_sender_volume$cum <- cumsum(top50_sender_volume$sender_volume)
1198 top50_sender_volume$perc <- top50_sender_volume$cum/tx_volume*100
plot(top50_sender_volume$perc, ylim = c(0,100),
1200   xlab = "Top 50 Senders",
   ylab = "Percentage of all Transactions Volume")
1202
1204
top50_receiver_volume <- head(m_w[order(-receiver_volume),c("wallet", "receiver_volume")], 50)
top50_receiver_volume$cum <- cumsum(top50_receiver_volume$receiver_volume)
1206 top50_receiver_volume$perc <- top50_receiver_volume$cum/tx_volume*100
plot(top50_receiver_volume$perc, ylim = c(0,100),
1208   xlab = "Top 50 Receivers",
   ylab = "Percentage of all Transactions Volume")
1210
plot(top50_sender_tx$perc, ylim = c(0,100), col = "blue",
1212   xlab = "Top 50 Senders (blue) and Top 50 Receivers (orange)",
   ylab = "Percentage of covered Transactions")
1214 points(top50_receiver_tx$perc, ylim = c(0,100), col = "orange")

1216 plot(top50_sender_volume$perc, ylim = c(0,100), col = "blue",
   # main = "Cumulative % of Volume covered by Major Users",
1218   xlab = "Top 50 Senders (blue) and Top 50 Receivers (orange)",
   ylab = "Percentage of covered Volume")
1220 points(top50_receiver_volume$perc, ylim = c(0,100), col = "orange")

```

```

1222 plot(top50_sender_tx$perc ,
        ylim = c(0,100) ,
1224         col = "blue" ,
        ylab = "Cumulative Percentage")
1226 points(top50_sender_volume$perc , ylim = c(0,100) , col = "orange")

1228 plot(top50_receiver_tx$perc , ylim = c(0,100) , col = "blue" ,
        ylab = "Cumulative Percentage")
1230 points(top50_receiver_volume$perc , ylim = c(0,100) , col = "orange")

1232
1234 intersect_volume <- intersect(top50_receiver_volume$wallet ,
                                top50_sender_volume$wallet)
1236 intersect_tx <- intersect(top50_receiver_tx$wallet ,
                             top50_sender_tx$wallet)

1238 intersect_senders <- intersect(top50_sender_tx$wallet ,
                                top50_sender_volume$wallet)
1240 intersect_receivers <- intersect(top50_receiver_tx$wallet ,
                                    top50_receiver_volume$wallet)

1242
1244 top50_wallet_tx <- head(m_w[order(-number_tx) , c("wallet" , "number_tx")], 50)
1246 top50_wallet_tx$cum <- cumsum(top50_wallet_tx$number_tx)
1248 top50_wallet_tx$perc <- top50_wallet_tx$cum/tx_n*100/2
1250 plot(top50_wallet_tx$perc , ylim = c(0,100) ,
        xlab = "Top 50 Wallets" ,
        ylab = "Percentage of all Transactions Number")

1252
1254 hist(m_w$balance/e_w_rate ,
        col = "orange" ,
        border = "grey" ,
        ylim = range(0,10) ,
        xlab = "Net Balances of Wallets , in ether" ,
        ylab = "Number of Wallets" , breaks = 50)

```

../analysis/data\_exploration.R

## REFERENCES

- [Albert et al., 2018] Albert, E., Gordillo, P., Livshits, B., Rubio, A., and Sergey, I. (2018). Ethir: A framework for high-level analysis of ethereum bytecode. In *International Symposium on Automated Technology for Verification and Analysis*, pages 513–520. Springer.
- [Alstott et al., 2014] Alstott, J., Bullmore, E., and Plenz, D. (2014). pow-erlaw: A python package for analysis of heavy-tailed distributions. *PLoS one*, 9:e85777.
- [Anderson et al., 2016] Anderson, L., Holz, R., Ponomarev, A., Rimba, P., and Weber, I. (2016). New kids on the block: an analysis of modern blockchains. *arXiv preprint arXiv:1606.06530*.
- [Androulaki et al., 2013] Androulaki, E., Karame, G. O., Roeschlin, M., Scherer, T., and Capkun, S. (2013). Evaluating user privacy in bitcoin. In *International Conference on Financial Cryptography and Data Security*, pages 34–51. Springer.
- [Anoica and Levard, 2018] Anoica, A. and Levard, H. (2018). Quantitative description of internal activity on the ethereum public blockchain. In *New Technologies, Mobility and Security (NTMS), 2018 9th IFIP International Conference on*, pages 1–5. IEEE.
- [Bartoletti et al., 2017] Bartoletti, M., Carta, S., Cimoli, T., and Saia, R. (2017). Dissecting ponzi schemes on ethereum: identification, analysis, and impact. *arXiv preprint arXiv:1703.03779*.
- [Baumann et al., 2014] Baumann, A., Fabian, B., and Lischke, M. (2014). Exploring the bitcoin network. volume 1.
- [Bennink et al., 2018] Bennink, P., Gijtenbeek, L. v., Deventer, O. v., and Everts, M. (2018). An analysis of atomic swaps on and between ethereum blockchains using smart contracts.
- [Bentke, 2017] Bentke, J. (2017). Analysis and realisation of acceptance criteria for pending transactions in an ethereum-blockchain environment.
- [Bistarelli et al., 2019] Bistarelli, S., Mazzante, G., Micheletti, M., Mostarda, L., and Tiezzi, F. (2019). Analysis of ethereum smart contracts and op-codes. In *International Conference on Advanced Information Networking and Applications*, pages 546–558. Springer.
- [BitInfoCharts, 2019a] BitInfoCharts (2015–2019a). Bitcoin (btc) price stats and information. [Online; accessed 10-03-2019].
- [BitInfoCharts, 2019b] BitInfoCharts (2015–2019b). Ethereum (eth) price stats and information. [Online; accessed 10-03-2019].
- [Buterik, 2013] Buterik, V. (2013). A next generation smart contract & decentralized application platform. [Online; accessed 2019-03-21].
- [Buterin et al., 2014] Buterin, V. et al. (2014). A next-generation smart contract and decentralized application platform. *white paper*.
- [Butin, 2015] Butin, V. (2015). Merkle in ethereum. online, access 14 Mar 2019.
- [Chan and Olmsted, 2017] Chan, W. and Olmsted, A. (2017). Ethereum transaction graph analysis. In *2017 12th International Conference for Internet Technology and Secured Transactions (ICITST)*, pages 498–500. IEEE.
- [Chandersekhar, 2018] Chandersekhar, P. (2018). Ethereum smart-contracts: Most of them are rarely used! online, access 14 Mar 2019.
- [Chen et al., 2018] Chen, T., Zhu, Y., Li, Z., Chen, J., Li, X., Luo, X., Lin, X., and Zhang, X. (2018). Understanding ethereum via graph analysis. In *IEEE INFOCOM 2018-IEEE Conference on Computer Communications*, pages 1484–1492. IEEE.
- [Dannen, 2017] Dannen, C. (2017). *Introducing Ethereum and Solidity: Foundations of Cryptocurrency and Blockchain Programming for Beginners*. Apress, Berkeley, CA, USA, 1st edition.
- [de Balthasar and Hernandez-Castro, 2017] de Balthasar, T. and Hernandez-Castro, J. (2017). An analysis of bitcoin laundry services.
- [Dennis and Disso, 2019] Dennis, R. and Disso, J. P. (2019). An analysis into the scalability of bitcoin and ethereum. In *Third International Congress on Information and Communication Technology*, pages 619–627. Springer.
- [Drainville, 2012] Drainville, D. (2012). An analysis of the bitcoin electronic cash system. *Waterloo, Canada: University of Waterloo*.
- [Grishchenko et al., 2018a] Grishchenko, I., Maffei, M., and Schneidewind, C. (2018a). Ethertrust: Sound static analysis of ethereum bytecode.
- [Grishchenko et al., 2018b] Grishchenko, I., Maffei, M., and Schneidewind, C. (2018b). Foundations and tools for the static analysis of ethereum smart contracts. In *International Conference on Computer Aided Verification*, pages 51–78. Springer.
- [Grishchenko et al., 2018c] Grishchenko, I., Maffei, M., and Schneidewind, C. (2018c). A semantic framework for the security analysis of ethereum smart contracts. In *International Conference on Principles of Security and Trust*, pages 243–269. Springer.
- [Hagberg et al., 2008] Hagberg, A. A., Schult, D. A., and Swart, P. J. (2008). Exploring network structure, dynamics, and function using networkx. In Varoquaux, G., Vaught, T., and Millman, J., editors, *Proceedings of the 7th Python in Science Conference*, pages 11 – 15, Pasadena, CA USA.
- [Kaminsky, 2011] Kaminsky, D. (2011). Black ops of tcp/ip. *Black Hat USA*, page 44.
- [Kiffer et al., 2017] Kiffer, L., Levin, D., and Mislove, A. (2017). Stick a fork in it: Analyzing the ethereum network partition. In *Proceedings of the 16th ACM Workshop on Hot Topics in Networks*, pages 94–100. ACM.
- [Kim et al., 2018] Kim, S. K., Ma, Z., Murali, S., Mason, J., Miller, A., and Bailey, M. (2018). Measuring ethereum network peers. In *Proceedings of the Internet Measurement Conference 2018*, pages 91–104. ACM.
- [Knorr, 2019] Knorr, C. (2019). *Holland/Leinhardt (1971): Transitivity in Structural Models of Small Groups*, pages 267–270.
- [Li et al., 2017] Li, Y., Zheng, K., Yan, Y., Liu, Q., and Zhou, X. (2017). Etherql: a query layer for blockchain system. In *International Conference on Database Systems for Advanced Applications*, pages 556–567. Springer.
- [Lischke and Fabian, 2016] Lischke, M. and Fabian, B. (2016). Analyzing the bitcoin network: The first four years. *Future Internet*, 8(1):7.
- [Maesa, 2018] Maesa, D. D. F. (2018). Blockchain applications: Bitcoin and beyond.
- [Meghanathan, 2018] Meghanathan, N. (2018). Graph theoretic approaches for analyzing large-scale social networks. pages 7–34. IGI Global.
- [Meiklejohn et al., 2013] Meiklejohn, S., Pomarole, M., Jordan, G., Levchenko, K., McCoy, D., Voelker, G. M., and Savage, S. (2013). A fistful of bitcoins: characterizing payments among men with no names. In *Proceedings of the 2013 conference on Internet measurement conference*, pages 127–140. ACM.
- [Newman, 2003] Newman, M. E. J. (2003). The structure and function of complex networks. *SIAM Review*, 45:167–256.
- [Ober et al., 2013] Ober, M., Katzenbeisser, S., and Hamacher, K. (2013). Structure and anonymity of the bitcoin transaction graph. *Future internet*, 5(2):237–250.
- [Ortega, 2013] Ortega, M. S. (2013). *The bitcoin transaction graphanonymity*. PhD thesis, Masters thesis, Universitat Oberta de Catalunya.
- [Partz, 2018] Partz, H. (2018). Ethereum unique addresses break 50 million, active wallet number keeps dropping - cointelegraph. [Online; accessed 10-03-2019].
- [Payette et al., 2017] Payette, J., Schwager, S., and Murphy, J. (2017). Characterizing the ethereum address space.
- [Pérez Jiménez, 2018] Pérez Jiménez, C. (2018). Analysis of the ethereum state.
- [Piper and Caver, 2018] Piper, M. and Caver, J. (2018). Web3py: A Python interface for interacting with the ethereum blockchain and ecosystem. [Online; accessed 10-03-2019].
- [Portal, 2019a] Portal, S. T. S. (2017–2019a). Price of ethereum from february 2017 to february 2019 (in u.s. dollars). [Online; accessed 10-03-2019].
- [Portal, 2019b] Portal, S. T. S. (2017–2019b). Size of the bitcoin blockchain from 2010 to 2019, by quarter (in megabytes). [Online; accessed 10-03-2019].
- [Pratama and Mutijarsa, 2018] Pratama, F. A. and Mutijarsa, K. (2018). Query support for data processing and analysis on ethereum blockchain. In *2018 International Symposium on Electronics and Smart Devices (ISESD)*, pages 1–5. IEEE.
- [Preethi, 2017] Preethi, K. (2017). [Online; accessed 2019-03-21].
- [Ray, 2018] Ray, J. (2018). Proof-of-stake faq’s. online, access 14 Mar 2019.
- [Reid and Harrigan, 2013] Reid, F. and Harrigan, M. (2013). An analysis of anonymity in the bitcoin system. In *Security and privacy in social networks*, pages 197–223. Springer.
- [Rossum, 1995] Rossum, G. (1995). Python reference manual. Technical report, Amsterdam, The Netherlands, The Netherlands.
- [Rudlang, 2017] Rudlang, M. (2017). Comparative analysis of bitcoin and ethereum. Master’s thesis, NTNU.
- [Sapuri et al., 2017] Sapuri, S., Kokkinaki, A., and Georgiou, I. (2017). In which distributed ledger do we trust? a comparative analysis of cryptocurrencies. *MCIS 2017 Proceedings*.
- [Signer, 2018] Signer, C. (2018). Gas cost analysis for ethereum smart contracts. Master’s thesis, ETH Zurich, Department of Computer Science.
- [Somin et al., 2018a] Somin, S., Gordon, G., and Altschuler, Y. (2018a). Network analysis of erc20 tokens trading on ethereum blockchain. In *International Conference on Complex Systems*, pages 439–450. Springer.

- [Somin et al., 2018b] Somin, S., Gordon, G., and Altshuler, Y. (2018b). Social signals in the ethereum trading network. *arXiv preprint arXiv:1805.12097*.
- [Spagnuolo et al., 2014] Spagnuolo, M., Maggi, F., and Zanero, S. (2014). Bitiodine: Extracting intelligence from the bitcoin network. In *International Conference on Financial Cryptography and Data Security*, pages 457–468. Springer.
- [Tikhomirov et al., 2018] Tikhomirov, S., Voskresenskaya, E., Ivanitskiy, I., Takhaviev, R., Marchenko, E., and Alexandrov, Y. (2018). Smartcheck: Static analysis of ethereum smart contracts.
- [Udaigiriya and Sharma, 2015] Udaigiriya, L. and Sharma, S. K. (2015). Complementary cumulative distribution function for performance analysis of ofdm signals and implement papr reduction techniques in ofdm system using ccdf function on matlab.
- [Watts and H. Strogatz, 1998] Watts, D. and H. Strogatz, S. (1998). Collective dynamics of small world networks. *Nature*, 393:440–2.
- [Wood, 2014] Wood, G. (2014). Ethereum: A secure decentralised generalised transaction ledger. *Ethereum project yellow paper*, 151:1–32.
- [Zwang et al., 2018] Zwang, M., Somin, S., Pentland, A., and Altshuler, Y. (2018). Detecting bot activity in the ethereum blockchain network. *arXiv preprint arXiv:1810.01591*.